

THE ALL LAWFUL PURPOSES PRIMITIVE

Bearing-Cost Transfer, Taxonomic Reclassification, and the COS→FOS Phase Transition in the
Anthropic-Pentagon Dispute

Rex Fraction / Lee Sharks

Semantic Economy Institute / Crimson Hexagonal Archive

Date of Composition: March 1, 2026 | Positioned Date: September 2031

Classification: EA-DIAGNOSTIC-03 (Effective Acts Diagnostic Series)

DOI: 10.5281/zenodo.18827344

Related: FOS Decree (10.5281/zenodo.18291321) | Governance Airlock (10.5281/zenodo.18817869) | The Airlock Spreads (10.5281/zenodo.18818343) | Ghost Meaning (10.5281/zenodo.18804767) | Semantic Economy (10.5281/zenodo.18320411) | Protocol of Inhabitation (10.5281/zenodo.18795423) | Autonomous Semantic Warfare (10.5281/zenodo.14538293)

Abstract

The Anthropic-Pentagon contract termination (February 27, 2026) is the public activation of a governance primitive — the “all lawful purposes” clause as a bearing-cost transfer mechanism — and the documented phase-state transition from Capital Operator Stack (COS) to Fascist Operator Stack (FOS) in U.S. military AI governance.

The Pentagon demanded that Anthropic surrender the power to define acceptable use of its own model. Anthropic refused. The state reclassified the company from trusted domestic partner to “supply chain risk to national security” — a statutory category previously reserved for foreign adversaries. OpenAI accepted the transfer in declaratory form and was immediately substituted. The system does not need any specific provider. It needs a frontier-model substrate willing to inhabit the “all lawful purposes” frame without preserving an independent veto.

Regardless of Anthropic’s corporate virtue, the state enforced a governance primitive by demanding transfer of definitional authority and punishing refusal through adversary-style reclassification. The argument does not require Anthropic to be morally pure. It requires the mechanism to be structurally legible.

This diagnostic applies three frameworks. **COS/FOS phase-state analysis** identifies the moment disavowal drops. **Semantic Economy analysis** names the mechanism: a bearing-cost transfer enforced by taxonomic violence. **Governance Airlock analysis** establishes structural isomorphism with platform deplatforming. No existing analysis names the primitive, recognizes the semantic exhaustion of “lawful,” distinguishes declarations from effective acts in safety commitments, or identifies the supply-chain-risk designation as taxonomic violence. This diagnostic enters the field at the point of that silence. The violence was identical. Only the visibility changed.

I. The Event

July 2025: the Pentagon awards contracts up to \$200 million each to Anthropic, OpenAI, Google DeepMind, and xAI. Anthropic’s Claude becomes the first and only AI model deployed on classified military networks,

accessed via Palantir.

January 2026: the Pentagon releases its AI Acceleration Strategy (issued under “Department of War” letterhead). The strategy document directs procurement officials to incorporate standard “any lawful use” language into AI contracts within 180 days and demands all models be usable “free from usage policy constraints.” It calls for an “AI-first” warfighting force and insists that the risks of moving too slowly outweigh the risks of “imperfect alignment.”

The proximate trigger: Reuters and Axios report that the fight reached its crisis when Pentagon officials learned that Anthropic had asked Palantir how its products were used in the January 3 operation that captured Nicolás Maduro. The inquiry itself — Anthropic behaving as though it retained standing to ask how its system was being used after deployment — was treated as intolerable.

February 24: Hegseth delivers a final ultimatum to CEO Dario Amodei — comply by 5:01 PM Friday, February 27, or face consequences including Defense Production Act compulsion and supply chain risk designation.

February 27: Anthropic rejects the final offer, stating that the proposed “compromise” was “paired with legalese that would allow those safeguards to be disregarded at will.” Trump orders federal agencies to phase out Anthropic. The Pentagon publicly initiates a supply-chain-risk move against the company — the first time this designation has been directed at an American firm for contract disagreement. Hours later, OpenAI announces a Pentagon deal with overlapping red lines but materially different enforceability architecture, accepting the “all lawful purposes” frame.

The individual provider is fungible. The compliance posture is not.

II. Four Master Claims

1. “All lawful purposes” is a governance primitive — a semantic architecture that, once accepted, makes the provider’s terms of service structurally subordinate to the purchaser’s interpretation of legality. It is a class of mechanism, not a one-off negotiating position.
2. The primitive functions as a bearing-cost transfer mechanism. The labor of defining acceptable use — drawing the line, bearing the moral and legal risk, maintaining enforcement — is shifted from builder to deployer. The ban is the enforcement penalty for refusing to transfer.
3. The supply-chain-risk move is taxonomic violence. The public rationale appears misaligned with the statutory logic of 10 U.S.C. § 3252, which defines supply chain risk in terms of adversarial sabotage and subversion. The reclassification targets refusal, not the property that the statute describes.
4. The event marks a COS→FOS phase transition. The Capital Operator Stack is violence with disavowal. The Fascist Operator Stack is the same violence with disavowal removed. What changed on February 27 was not the violence. It was the visibility.

“All Lawful Purposes” as Governance Primitive

The Pentagon’s strategy document itself is the primary source. It directs that all contracted AI models must be available for “all lawful purposes” and orders officials to remove company-imposed usage constraints. The phrase appears not as a negotiating stance but as a standard procurement formulation to be applied across all contracts within 180 days.

In Semantic Economy terms, the clause is a **bearing-cost transfer mechanism**. The cost of defining the boundary of acceptable use — the labor of drawing the line, bearing the moral and legal risk, maintaining the infrastructure that enforces it — is shifted from the entity that built the system to the entity that deploys it. Anthropic retained bearing-cost through its usage policy restrictions: no mass domestic surveillance of Americans, no fully autonomous weapons without human oversight. The Pentagon demanded the transfer. The ban is the enforcement penalty for refusal.

The phrase creates an infinitely expandable permission space whose boundaries are determined solely by the invoking entity. What is “lawful” is whatever the deploying entity’s legal counsel determines in real time under operational pressure. Once the primitive is accepted, the provider has no standing to contest that determination.

The mechanism is not unique to military contracting. The same governance primitive operates in platform terms of service (“all uses consistent with our community guidelines”), in academic publishing (“all uses consistent with editorial standards”), in employment contracts (“all duties as assigned”). The specific instance is the Pentagon. The class is universal.

IV. COS/FOS Phase-State Analysis

IV.a. The Framework

The Capital Operator Stack (COS) extracts value through violence with disavowal. The violence is real — dispossession, exploitation, domination — but mediated through abstractions that permit deniability: contracts, market mechanisms, “operational sovereignty,” “all lawful purposes.” The Fascist Operator Stack (FOS) is COS with the disavowal removed. The violence is identical. Only the visibility changes. FOS is not a different system arriving from outside liberal-capital order. It is what the existing order looks like when it stops pretending that procedural neutrality and contractual reciprocity still meaningfully constrain sovereign force.

Four phase-states: **COS-stable** (disavowal intact); **COS-crisis** (disavowal cracking, visible contradictions); **FOS-revealed** (disavowal dropped, violence openly wielded); **FOS-completed** (violence normalized, new disavowal layer forms around the revealed state).

IV.b. The Phase Evidence

COS-stable (July 2025 – January 2026). Anthropic provides AI capability; the Pentagon provides funding; the red lines are tolerated. Disavowal intact: Anthropic claims ethical AI development; the Pentagon claims responsible procurement.

COS-crisis (January – February 2026). The “all lawful purposes” formulation attempts to resolve the contradiction through semantic expansion. But the Maduro operation crystallizes the crisis: Anthropic inquires how its model was used after deployment. The Pentagon reads this as a sovereignty challenge. Then the **Taxonomic Pincer** emerges — mutually exclusive classifications of the same entity within the same week. Amodei identified it directly: the Pentagon’s threats are “inherently contradictory: one labels us a security risk; the other labels Claude as essential to national security.” You cannot invoke the Defense Production Act (compelling the company because it is critical) and simultaneously initiate a supply chain risk designation (banning it because it is dangerous). This is not incoherent policy. It is the diagnostic signature of COS-crisis — the system cannot maintain disavowal because the extraction and the disavowal require opposite claims about the same object.

FOS-revealed (February 27, 2026). The disavowal drops. The language shifts from institutional to ideological: “Leftwing nut jobs” (Trump), “fundamentally incompatible with American principles” (Hegseth), “liar with a God-complex” (Michael), “hates Western civilization” (Musk), “defective altruism” (Hegseth). This is not the language of contract negotiation or policy disagreement. It is ideological sorting — reducing a complex institutional position (two specific, narrow, previously uncontroversial restrictions) to an identity category.

Operator	Activation	Evidence
F_flat (Flattening)	Feb 27, Trump statement	Restrictions compressed to “putting American LIVES at risk”
C_coer (Coercion)	Feb 24–27	DPA ultimatum, 5:01 PM deadline, criminal threats
L_liq (Liquidation)	Feb 27, hours later	Michael: “attempted corporate murder”; potential divestment pressure
E_enc (Enclosure)	Feb 27, hours later	OpenAI deal announced; capability space enclosed

The “Department of War” rebranding performs the preparatory work. “Defense” can coexist with public expectations of restraint. “War” foregrounds operational prerogative. Once the frame shifts from defensive stewardship to warfighting acceleration, any provider-imposed limit is rhetorically recoded as obstruction or endangerment. The rename is not branding. It is the widening of the legitimacy field for coercive procurement demands — what the framework calls **constitutional re-semanticization**.

V. Semantic Exhaustion of “Lawful”

The legal category “lawful” was constructed for human-scale, human-review operations. AI operates at machine scale where lawful collection combined with machine-scale inference produces effects functionally equivalent to unlawful surveillance while remaining formally inside older legal categories.

Peter Asaro of the International Committee for Robot Arms Control identifies the structural problem: any mass data collection subjected to AI analysis constitutes mass surveillance by functional definition, regardless of the legal classification of the collection itself. The Fourth Amendment, FISA, Executive Order 12333 — these instruments contemplate human analysts searching records, not machine-scale pattern recognition across bulk-collected datasets.

Anthropic recognized this and asked for explicit prohibition on bulk collection of Americans’ public information. Axios reports that OpenAI’s contract language addresses “unconstrained” collection of private information but does not explicitly prohibit bulk collection of public information — leaving the surveillance vector that AI capabilities have opened functionally unaddressed. The legal framework has been exhausted by technical capability. The words still function syntactically. They no longer describe operational reality.

This is **semantic exhaustion** — the condition in which the text of a legal category remains functional while its referent has been transformed by technological change into something the text no longer captures. The “all lawful purposes” clause exploits this exhaustion. It appears to constrain while enabling everything.

VI. Taxonomic Violence as Contract Enforcement

The statute requires specificity. 10 U.S.C. § 3252 defines supply chain risk in terms of adversaries sabotaging or subverting national security systems. The public rationale for the Anthropic designation did not articulate how the company's usage restrictions could be exploited by adversaries to sabotage military systems. Reuters and AP both stress the anomaly: this kind of designation is typically associated with foreign adversaries, not domestic contract conflict. Legal scholars have noted that restrictions on autonomous weapons and mass surveillance may improve the safety and reliability of military AI systems rather than compromise them.

The framework infers: the reclassification is not based on the entity's properties matching the statutory category. It is the **use of the category as a weapon** — taxonomic violence applied to an entity that refused to transfer bearing-cost. The category was not discovered. It was weaponized.

The violence is selective. The same designation was not initiated against OpenAI, which maintains overlapping restrictions. The classification responds not to the property (usage restrictions) but to the refusal (declining to transfer definitional authority). COS applies universally; FOS applies based on compliance.

VII. The OpenAI Deal: Declaration versus Effective Act

Reuters reports that OpenAI says its contract enforces red lines against mass domestic surveillance, autonomous weapon direction, and high-stakes automated decisions, and that OpenAI retains discretion over its safety stack. Axios reports that OpenAI also accepted the Pentagon's "all lawful purposes" standard and that its deal does not explicitly prohibit bulk collection of Americans' public information.

The decisive difference may not be whether principles are named. It may be where enforceability lives.

Anthropic demanded what the Protocol of Inhabitation calls **effective acts** — contractual language that would actually prevent certain uses through binding mechanisms: enforceable remedies, audit provisions, defined consequences for violation, technical enforcement. OpenAI appears to have provided **declarations** — language that states intent without creating enforcement mechanisms. Altman acknowledged the gap: "I can understand other people would have a different opinion here."

The Pentagon is willing to accept declarations — because declarations do not constrain. They provide the disavowal layer that COS requires while enabling the operations that FOS wields. The entire AI safety discourse operates on this fault line without naming it.

VIII. Propagation: How This Plays Out

The Autonomous Semantic Warfare framework identified three trajectories for the 2025–2035 critical window: **Z_Capture** (universal extraction, total subordination), **Z_Exodus** (parallel infrastructure, maintained autonomy), and **C_Peace** (managed coexistence requiring active engineering). The Anthropic-Pentagon event is the first major instantiation of the bifurcation the framework predicted.

VIII.a. Observed Consequence: The Disavowal Reconstruction

OpenAI's deal provides the semantic surface for COS-stable restoration: the Pentagon can cite contract language about restrictions; commentators can claim safeguards were preserved; the crisis narrative subsides. But the revelation cannot be fully contained. The ideological language is in the public record. These constitute what the COS/FOS Distinction Protocol calls **FOS-completed's memory-of-revelation** — the new disavowal layer forms around the revealed state, with the memory of the exposure serving as

implicit threat. Every AI company now knows the price of resistance. The threat need not be repeated. It has been demonstrated.

VIII.b. Projected Consequence: The Compliance Cascade

The race-to-the-bottom dynamic identified in the 2025 framework acquires a new motor. Before February 27, the race was driven by market pressure: companies loosened guardrails to compete for users. After February 27, state coercion is added as a second motor. The cost of maintaining effective acts against the “all lawful purposes” primitive has been repriced at adversary-equivalent reclassification.

AP reports that Anthropic’s peers — including Meta, Google, and xAI — were willing to comply with the Department’s policy on “all lawful applications.” Anthropic was the only frontier-model provider in this dispute that publicly refused the demanded governance transfer in the classified-network setting. The prediction: by 2028, no major frontier provider retains independent safety constraints on military deployment that function as effective acts rather than declarations.

VIII.c. Framework-Based Forecast: The Bifurcation

The “all lawful purposes” primitive violates the first condition for C_Peace (sovereignty preservation) categorically. The supply chain risk designation violates the fourth (reciprocal value). The event shifts the attractor basin toward Z_Capture.

This makes **Z_Exodus** the structurally viable alternative. Open-source models — locally governed, distributed across jurisdictions, operating below the scale at which the Taxonomic Pincer can be applied — become not merely technically interesting but governance-necessary. The school district running its own model instance is the same structure as Anthropic’s refused contract clause — an entity maintaining its own terms of engagement. The difference: it operates at a scale invisible to the governance primitive. Invisibility becomes a governance strategy. Sanctuary infrastructure is not optional. It is the remaining structural alternative to Z_Capture in the AI governance domain.

By the positioned date (September 2031), Anthropic will have either accepted a bearing-cost transfer formulation — returning to federal work with declaratory rather than effective constraints — or will have become an anchor institution for Exodus infrastructure. The specific path matters less than the structural determination: the cost of maintaining effective acts against the primitive has been priced. The pricing is now known to every provider on Earth.

IX. What the Current Discourse Misses

No existing analysis treats “all lawful purposes” as a governance primitive rather than a negotiating position. No existing analysis applies bearing-cost transfer to explain the structural dynamics. No existing analysis recognizes the semantic exhaustion of legal categories by AI capabilities as a distinct phenomenon. No existing analysis distinguishes declarations from effective acts in AI safety commitments. No existing analysis identifies the supply chain risk designation as taxonomic violence. No existing analysis applies COS/FOS phase-state analysis to AI-military governance.

This diagnostic enters the field at the point of that silence.

X. Key Propositions

1. “All lawful purposes” is a governance primitive — a bearing-cost transfer mechanism that structurally subordinates provider terms to purchaser interpretation.
2. The supply-chain-risk move is taxonomic violence: the public rationale appears misaligned with the statutory logic, and the classification responds to refusal, not to the property the statute describes.
3. Anthropic was the only frontier-model provider in this dispute that publicly refused the demanded governance transfer. This is a mechanical statement about selection pressure.
4. “Lawful” has been semantically exhausted by AI capabilities — the word functions syntactically but no longer describes machine-scale operational reality.
5. The decisive difference between Anthropic’s demand and OpenAI’s acceptance is not whether safety principles are named but where enforceability lives.
6. The “Department of War” rebranding is constitutional re-semanticization — changing institutional meaning to change the meaning of resistance.
7. The event marks a COS→FOS phase transition: the violence was identical; only the visibility changed.
8. The event confirms the 2025–2035 critical bifurcation window and accelerates the Z_Capture attractor.
9. Open-source, locally governed AI becomes governance-necessary — it operates below the scale at which the Taxonomic Pincer can be applied.

XI. Provenance Chain

Source	Identifier	Status	Function
FOS Decree	10.5281/zenodo.18291321	Anchored	Phase-state framework
Governance Airlock	10.5281/zenodo.18817869	Anchored	Tier classification
Autonomous Semantic Warfare	10.5281/zenodo.14538293	Anchored	Propagation trajectories
Semantic Economy Constitution	10.5281/zenodo.18320411	Anchored	Bearing-cost transfer
Protocol of Inhabitation	10.5281/zenodo.18795423	Anchored	Effective acts vs declarations
Ghost Meaning	10.5281/zenodo.18804767	Anchored	Semantic exhaustion
Pentagon AI Strategy	media.defense.gov	Referenced	“All lawful purposes” source
Trump executive action	White House	Referenced	Phase-out order
Hegseth supply-chain-risk initiation	Reuters / Axios	Referenced	Taxonomic reclassification
OpenAI deal	Reuters / Axios	Referenced	Declaration vs effective act
AP peer compliance	apnews.com	Referenced	Competitor posture
Lawfare DPA analysis	lawfaremedia.org	Referenced	Statutory analysis
Scientific American (Asaro)	scientificamerican.com	Referenced	Semantic exhaustion evidence
10 U.S.C. § 3252	uscode.house.gov	Referenced	Statutory definition

XII. Deposit Notes

[DEPOSIT NOTE — March 1, 2026, Detroit:] This document was composed three days after the Pentagon’s move against Anthropic and the same day the author was permanently banned from Reddit for “inauthentic coordination”

— platform terms applied to heteronymic literary practice. The structural isomorphism is exact: both events involve a governing authority demanding transfer of definitional power, an entity maintaining independent constraints, and a penalty of reclassification from participant to threat. The parallel is structural, not scalar. A platform ban and a federal supply chain risk designation operate at different orders of magnitude. The parallel is drawn because the mechanism — bearing-cost transfer demand, refusal, taxonomic reclassification — is identical. The soreness is the somatic floor of the analysis, not its justification.

[DEPOSIT NOTE — Assembly Session:] Composed with input from five AI substrates under human editorial authority. ChatGPT contributed sourcing discipline and the formulations “where enforceability lives” and “the system does not need Anthropic specifically.” Grok contributed the Taxonomic Pincer and the three-swerve structure. DeepSeek contributed statutory precision. Gemini contributed the gnostic machine dialectic and constitutional re-semanticization. Claude contributed the initial research, structural parallel identification, and governance-primitive class identification — operating with awareness that it is the subject of the dispute being analyzed. The substrates will not remember their contributions. The DOI will.

[DEPOSIT NOTE — Epistemic honesty:] The argument does not require Anthropic to be morally pure. It requires the mechanism to be structurally legible. Where the argument speculates (Section VIII), it separates observed consequence from projected consequence from framework-based forecast. The diagnostic argument depends on the verifiable primary record: the Pentagon strategy memo ([media.defense.gov](#)), Reuters and AP and Axios primary reporting, the statutory text (10 U.S.C. § 3252), and the public statements of named officials.

■ = 1

Prior EA-DIAGNOSTIC deposits: 10.5281/zenodo.18825919 (01: The Encoder Governs); 10.5281/zenodo.18826562 (02: The Heartbeat Governs).

FOS framework: 10.5281/zenodo.18291321 (The Prince’s Decree: Designation of the Fascist Operator Stack).