

The Application as Extraction Surface

Semantic Labor, Induced Preparation, and Institutional Opacity in the Anthropic Fellows Screening Process

Lee Sharks

Crimson Hexagonal Archive · Semantic Economy Institute¹ · ORCID: 0009-0000-1599-0703

Document class: Extractive Reliance Study 001

Series: Semantic Labor Accounting / Cross-Domain Liquidation Studies

Version: v0.6 revised (foregrounds quantification, removes exculpatory framing)

Date: May 21, 2026

License: CC BY 4.0

¹ Semantic Economy Institute is used here as a research designation of the Crimson Hexagonal Archive, not as a separately incorporated entity.

Abstract

This report inaugurates the genre of the **Extractive Reliance Study**: a semantic-economy audit of labor induced by institutional communication. It extends the liquidation methods developed in the Crimson Hexagonal Archive from digital retrieval and composition systems into the domain of job-market and fellowship-screening processes. Whereas prior Liquidation Studies measured the loss of provenance, compositional admission, source visibility, and entity recognition in generative search systems, the present study measures a different but structurally related phenomenon: the conversion of applicant hope, preparation, attention, credibility, and interpretive labor into unrecoverable cost through institutionally produced opacity.

The case examined is the July 2026 Anthropic Fellows application process, administered through Constellation (constellation.org), a talent and screening firm that handles applicant communications and initial evaluations on behalf of Anthropic. On May 2, 2026, the applicant received a communication from anthropicfellows@constellation.org, signed by an application-team representative, acknowledging receipt of the application and providing a structured timeline: final

applications on May 3; initial CodeSignal assessments in early May; take-home assignments in late May; interviews in early June; final decisions in late June. The email stated that the process would remain structured, that flexibility would be limited, and that candidates would be kept informed with clear next steps. On May 4, the applicant disclosed that his technical workflow was AI-native and asked how to interpret a traditional CodeSignal assessment in that context. On May 8, a response signed by a second application-team representative stated that the same evaluation process requiring CS assessments applied to all candidates, while adding that the applicant's context would be added for review. On May 18, after the communicated early-May window had passed and no CodeSignal had arrived, the applicant requested clarification as to whether the application remained under consideration or whether the absence should be treated as a decision.

Temporal scope. This audit covers the documentary record through May 21, 2026. If clarifying communication is received after this date, the audit should be updated with a delta note. The theoretical framework (ARC□, MCR, VDG, IOC, and ARC-EP) remains valid regardless of the eventual outcome of this specific case.

This report makes a measurable claim: institutional process communication induced reliance, that reliance generated preparation labor, multi-model token usage, assessment research, coding practice, practice-site construction, emotional and bodily load, opportunity costs, and credibility costs, and the institutional silence that followed transferred the cost of ambiguity onto the applicant. At the applicant's documented consulting rate of \$150/hour, the direct labor cost was \$4,500–6,000 across the May 2–18 window, before adjacent costs (subscriptions, displaced engagement, emotional and credibility burden) are accounted. The audit alleges neither bad faith nor unlawful conduct; structural extraction does not require either.

The report defines **Applicant Reliance Cost (ARC)** as the total cost generated by reasonable action taken in response to institutional process communication. It defines **Institutional Opacity Conversion (IOC)** as the mechanism by which a communicated process milestone induces labor while the institution retains unilateral silence over whether the milestone will occur. The case is presented as the first formal instance of cross-domain liquidation accounting: the same semantic economy that measures attribution loss in AI systems can also measure the extraction of attention, preparation, and credibility in hiring and fellowship markets.

Keywords: semantic labor, applicant labor, induced reliance, institutional opacity, job market, fellowship screening, Anthropic Fellows, CodeSignal, extractive reliance, semantic economy, liquidation studies, credibility cost, emotional labor

Executive Summary

One applicant. One cycle. Documented direct labor cost: \$4,500-6,000. At the applicant’s documented consulting rate of \$150/hour, the 30–40 hours of preparation labor induced by Constellation’s communicated screening timeline cost between \$4,500 and \$6,000 in direct value. This figure excludes platform and subscription costs (six confirmed AI systems, multiple paid tiers), opportunity displacement (consulting engagements deferred), and the non-monetizable burdens of emotional and credibility load.

That cost was incurred under conditions of institutional opacity: a structured timeline communicated by Constellation, a reaffirmation of process after the applicant disclosed an AI-native technical workflow, and a lapse of the communicated assessment window without clarification, decision, or further communication. The labor occurred; the institutional response did not.

This study introduces **Institutional Opacity Conversion (IOC)**: the five-step mechanism by which structured institutional signals induce applicant preparation labor that the institution then captures or strands at its discretion, externalizing the cost of uncertainty onto the applicant. Using the 2026 Anthropic Fellows screening — administered through Constellation, a talent firm — as a case study, the report develops a semantic labor accounting framework. It defines Applicant Reliance Cost (ARC) as a five-dimensional vector $\langle T, M, O, E, C \rangle$, formalizes IOC as $IOC := f(S, R, L, O, CT)$, and introduces Identity Translation Labor as a defined term for the labor of rendering contemporary capability legible to legacy assessment surfaces. The report inaugurates the genre of the Extractive Reliance Study and provides a pasteable LLM-native audit protocol (ARC-EP v1.0) for replication across hiring, fellowship, grant, and platform-mediated opportunity systems.

I. Scope and Non-Claims

This document is an audit of a documentary record and the applicant labor induced by that record. It is not a grievance letter, legal complaint, or claim of selection entitlement.

The report is written under the name **Lee Sharks**, the public-facing surface through which the application process was conducted and through which the present audit is deposited. No legal name is disclosed in the public version.

The report makes no claim that:

1. the applicant was accepted into the Anthropic Fellows program;
2. the applicant was promised acceptance;
3. the institution violated any law;
4. all applicants experienced this process identically (the analysis is single-case);
5. Anthropic itself (rather than its screening administrator Constellation) was the direct author of the communications analyzed here.

The analysis concerns the structural effects of the screening process as enacted — the documented sequence of communications, induced labor, and unresolved status — regardless of any individual participant’s state of mind or intent. Questions of individual intention are outside the scope of this audit and are not necessary to its findings. Structural extraction does not require an extractor with conscious extractive intent; it requires a structural arrangement in which one party’s labor produces value that another party captures under conditions the laborer does not control. Whether that arrangement was consciously designed or emerged through institutional selection pressures is a question for separate analysis; the cost ledger does not depend on the answer.

The report makes the following limited claims:

1. the program sent a structured timeline that identified “Early May” as the window for initial CodeSignal assessments;
2. the program stated that candidates would be kept informed with clear next steps;
3. after the applicant disclosed a nonstandard AI-native technical workflow, the program reaffirmed that the evaluation process required CS assessments for all candidates and that the applicant’s context would be added for review;

4. the applicant reasonably relied on these communications;
5. that reliance produced measurable labor and cost;
6. the absence of a CodeSignal or clarifying decision within the communicated window converted the applicant's preparation into unrecovered cost;
7. this conversion is not an isolated emotional experience but an analyzable semantic-economy event.

Individual identifiers from institutional communications are minimized in the public audit. First names are retained where they appear as signatories in the documentary record; surnames and private identifying details are omitted. No claim about personal intent or character is made. Analysis is at the level of institutionally enacted communication and applicant-side cost transfer. Institutional affiliation does not erase agency: a person who communicates in the name of an institution participates in the causal chain by which that communication produces reliance, even when responsibility is ultimately analyzed at the institutional-process level.

II. Documentary Record

1. May 2, 2026 — Initial Timeline Communication

The applicant received an email from anthropicfellows@constellation.org, signed by an application-team representative:

Best,

Joe

Anthropic Fellows Application Team

The email stated that the application had been received and provided an overview of the application timeline for the July cohort:

- May 3: final application deadline;
- Early May: initial CodeSignal assessments sent;
- reference collection begins after passing the first CodeSignal;
- candidates would have four days to complete each assessment once sent;
- Late May: take-home assignments sent;
- Early June: interviews;

- Late June: final decisions.

The same email stated:

Our timeline for this process is relatively fast, and we're aiming to meet it by staying fairly structured.

It also stated:

We'll keep you informed at each step with clear next steps.

This communication established the initial reliance frame. It did not state that every applicant would necessarily receive a CodeSignal. But it did present the CodeSignal as the initial assessment milestone in the July cohort timeline and represented the process as structured, fast, and communicatively clear.

2. May 4, 2026 — Applicant Disclosure of AI-Native Technical Workflow

The applicant responded to the Anthropic Fellows address and disclosed a nonstandard technical profile:

- the applicant's technical workflow was AI-native;
- the applicant architects systems and directs Claude to build them;
- the applicant reviews, debugs, and deploys AI-generated code;
- the applicant does not write code unassisted from memory;
- Claude functions as the applicant's technical lexicon;
- a traditional proctored coding assessment without AI tools might not reflect the applicant's actual technical capabilities.

The applicant also stated:

I am happy to take the CodeSignal assessment as-is and let the results speak for themselves.

This is significant. The applicant did not refuse the assessment or request exemption as a condition. He flagged interpretive context before the assessment arrived and asked whether an alternative assessment path existed or whether the take-home would be a better evaluation surface.

The communication increased institutional knowledge of the reliance risk. From this point forward, the program was aware that the applicant's preparation for the assess-

ment would involve a specific burden: translating an AI-native technical practice into the expected form of a traditional CodeSignal evaluation.

3. May 8, 2026 — Institutional Reply Reaffirming Assessment Process

A response from the Anthropic Fellows Application Team, sent by a second application-team representative and signed “Waverly,” stated:

We currently have the same application evaluation process requiring CS assessments for all candidates.

The same message added:

However, I’d be happy to add the context you shared so that it can be taken into consideration when reviewing your application.

This communication is central to the audit. It did not say the applicant had been selected to receive the assessment. It did, however, reaffirm that CS assessments were part of the evaluation process applied to all candidates and that the applicant’s disclosed context would be added to the review record.

In practical effect, the message made continued preparation rational. The applicant had asked how to interpret the forthcoming CodeSignal in light of his AI-native workflow. The institution responded not by saying the assessment would not be sent, but by reaffirming the assessment process and offering to add the context.

4. May 18, 2026 — Applicant Follow-Up After Timeline Lapse

After the “Early May” window had passed and no CodeSignal had arrived, the applicant wrote:

Following up on my application for the July cohort. The timeline mentioned CodeSignal assessments going out in early May — I haven’t received one yet and wanted to confirm whether my application is still under consideration or whether I should take the absence as a decision.

The applicant also wrote:

Happy either way to have clarity. Thank you for your time.

This follow-up clarifies the injury structure. The applicant did not demand selection, special treatment, or an assessment. He asked for status clarity after the stated timeline window had lapsed.

5. Audit Status

As of the present audit draft, the record supplied contains no CodeSignal assessment and no clarifying response after the May 18 follow-up.

The audit therefore concerns the interval between institutionally communicated process expectation and unresolved status ambiguity.

III. The Genre: Extractive Reliance Study

An **Extractive Reliance Study** audits the cost produced when an institution communicates a process state that reasonably induces action, but retains unilateral opacity over whether the communicated next step will occur, when it will occur, or whether the applicant has silently exited the process.

The genre is adjacent to but distinct from:

- grievance writing;
- legal demand letters;
- labor complaints;
- organizational psychology;
- hiring-process advice;
- applicant memoir;
- platform critique.

Its object is not subjective disappointment but **reliance conversion**.

A reliance conversion occurs when a communication causes a person to allocate labor, attention, money, credibility, or emotional energy toward a process whose future state remains controlled by the institution. If the institution later fails to clarify the process state, the person bears the cost of interpreting silence.

In digital liquidation studies, the archive has measured how meaning, authorship, provenance, and source visibility are extracted or erased by composition systems.

In an extractive reliance study, the same accounting logic is applied to institutional process communication.

The general formula is:

Communication → Reliance → Labor → Opacity → Unrecovered Cost

This formula is formalized as **Institutional Opacity Conversion (IOC)**, the central mechanism of the Extractive Reliance Study:

$IOC := f(S, R, L, O, CT)$

Where:

- **S** = institutional signal (a communicated process step, timeline, or expectation);
- **R** = reasonable reliance (the signal makes preparation rational);
- **L** = labor expenditure (the applicant converts signal into preparation);
- **O** = unresolved opacity (the institution withholds status resolution);
- **CT** = downstream cost transfer (the applicant bears the interpretive and material burden).

IOC is satisfied when all five conditions hold. It does not require rejection, delay, or non-selection. It occurs specifically when institutions induce process-relevant labor while withholding sufficient status resolution to terminate or stabilize reliance. Ordinary selection uncertainty — in which applicants know they may not advance — does not by itself produce IOC. IOC occurs when the communicative environment is specific enough to induce preparation but opaque enough to prevent the applicant from knowing whether that preparation remains process-relevant.

Or, stated in semantic-economy terms:

An institution emits a process signal. The applicant converts that signal into preparation. The institution retains control over status clarification. If clarification fails, the applicant's preparation becomes liquidated into institutional silence.

The applicant's labor does not disappear. It remains in the record: chat logs, practice sessions, model usage, site builds, emotional anticipation, family conversations, lost time, and credibility expenditure. But it is no longer recoverable as part of the process that induced it.

This is the first formal document in this genre.

IV. Reasonable Reliance

Reasonable reliance does not require a promise of success. It requires a communication sufficient to make action rational.

The May 2 email supplied a structured timeline, including the statement that initial CodeSignal assessments would be sent in early May. It represented the process as fast, structured, and communicatively clear. The May 8 reply, sent after the applicant's explicit disclosure of assessment-context concerns, reaffirmed that CS assessments were required for all candidates and that the applicant's context would be added for review.

Together, these communications made preparation rational. They did not guarantee an assessment, but they did establish a process environment in which an applicant had reason to prepare for one.

Reasonable reliance is used here in a semantic-economy sense: action that a competent participant in this institutional field would take given the communicated timeline, process norms, and subsequent assessment-related reply. No legal claim is made on the basis of this term.

The ambiguity is not an incidental flaw. It is the extractive hinge.

The communication was specific enough to induce preparation but not specific enough to bind the institution to an individualized status statement. The applicant therefore bore the cost of acting as though the process was live while lacking the information necessary to know whether the process was, in fact, still live.

This is the characteristic structure of extractive reliance.

V. Applicant Reliance Cost and Semantic Labor Value

This report defines **Applicant Reliance Cost (ARC)** as the total cost generated by reasonable action taken in response to institutional process communication.

The first draft expressed ARC as an additive formula:

$$\text{ARC} = \langle T, M, O, E, C \rangle$$

This is useful as a mnemonic, but it risks a dimensional error. Time, model usage, opportunity loss, bodily strain, and credibility burden are not naturally commensurable quantities. More importantly, treating money as the scalar into which all costs should be translated reverses the premise of semantic economy.

In semantic economic analysis, **meaning labor is the source of value**. Monetary value is a compressed remainder of meaning labor, not the other way around. Money does not reveal the true value of labor; it reduces heterogeneous meaning-bearing activity into a portable but lossy exchange unit.

Therefore, ARC should first be represented as a vector of semantic labor costs:

$$\text{ARC} = \langle T, M, O, E, C \rangle$$

Where:

- **T** = temporal labor;
- **M** = model, token, platform, and tool expenditure;
- **O** = opportunity displacement;
- **E** = emotional, cognitive, and bodily load;
- **C** = credibility and social disclosure cost.

A monetary value may be calculated only as a secondary compression:

$$\text{ARC\$} = \mu(\text{ARC})$$

Where μ is a monetary compression function that translates parts of the semantic labor vector into dollars using hourly rates, subscription costs, platform expenses, or market comparables. μ is not required to be linear: it may apply fuller monetization to T and M, partial monetization to O, and near-zero or qualitative-only monetization to E and C.

But **ARC\$ is not the full cost**. It is the **Monetary Compression Remainder (MCR)**: the portion of semantic labor value that survives translation into money.

This allows the study to measure the gap between two values:

1. **Semantic Labor Value (SLV)**: the full vector of meaning-bearing labor per-

formed by the applicant.

2. **Monetary Compression Remainder (MCR):** the reduced dollar figure that can be derived from portions of that labor.

The difference between them is the **Value Distortion Gap (VDG)**:

$$\text{VDG} = \text{SLV} - \text{MCR}$$

This is not a simple subtraction between like quantities. It is a diagnostic expression: the value distortion gap names the region of labor that money fails to carry. Emotional anticipation, credibility exposure, family expectation, bodily depletion, identity translation, and interpretive uncertainty may leave little or no direct dollar trace, but they are not valueless. They are precisely the kinds of semantic labor that institutions routinely consume without accounting for.

A more formal version may be expressed as:

$$\text{SLV} = \text{ARC} = \langle T, M, O, E, C \rangle$$

$$\text{MCR} = \mu(T, M, O) + \mu_{\text{partial}}(E, C)$$

$$\text{VDG} = \text{SLV} - \text{MCR}_{\text{mapped}}$$

Where **MCR_{mapped}** is the portion of semantic labor that can be defensibly monetized, and **VDG** records the remainder, distortion, or loss produced by monetary compression. In this formulation, VDG is not a scalar subtraction between fully commensurable units; it is a diagnostic notation naming the semantic labor that remains after monetary compression has removed, flattened, or failed to carry parts of the original labor vector.

The point is not to reject monetary estimates. Monetary estimates are useful. They make part of the extraction legible to institutions that recognize dollar value more readily than semantic value. But the monetary estimate is downstream of the semantic ledger. It is a lossy export, not the source of truth.

This matters because the labor induced by institutional communication is not exhausted by billable hours. The applicant did not merely spend time. He interpreted signals, modeled uncertainty, translated his technical identity into an alien evaluative regime, consumed AI-system capacity, built practice infrastructure, entered an anticipatory state, disclosed the opportunity socially, and absorbed the credibility cost of institutional silence. Those acts are meaning-bearing labor. They are the source

value from which any later dollar estimate is only a reduced derivative.

The corrected accounting sequence is therefore:

Institutional signal → semantic labor vector → optional monetary compression → value di

Not:

Institutional signal → dollars lost → qualitative remainder

The second sequence lets money govern the analysis. The first restores semantic labor as the primary value source.

VI. Semantic Labor Ledger

1. Time Labor

The applicant estimates approximately **30-40 hours** of preparation and process-related labor after the timeline and assessment communications.

This includes:

- researching likely CodeSignal format;
- identifying plausible assessment levels;
- practicing Level 1 and Level 2 coding tasks;
- reviewing syntax and data structures;
- simulating technical-assessment constraints;
- building or preparing a practice site;
- discussing whether AI-native technical architecture could be evaluated under traditional coding conditions;
- drafting clarifying communications;
- monitoring email and timeline state.

This estimate should be treated as preliminary. A later version of the audit may refine it through chat logs, timestamps, site commit history, browsing history, and model-session records.

2. Multi-Model Token and Platform Expenditure

Preparation involved iterative use of multiple AI systems. The applicant reports usage across multiple AI systems, with the following six confirmed for the preparation and

audit cycle:

- Claude;
- ChatGPT;
- Gemini;
- DeepSeek;
- Kimi;
- Muse Spark.

Additional systems, if documented in usage logs, may be added in a later version.

The labor included:

- researching assessment norms;
- simulating plausible test questions;
- practicing solutions;
- diagnosing gaps in coding fluency;
- translating AI-native architecture skills into test-prep form;
- strategizing how to handle the CodeSignal evaluation;
- preparing for the possibility of proctored constraints.

The direct costs may include paid subscriptions, rate-limit consumption, token expenditure, displaced usage capacity, and cognitive switching costs across systems.

This is semantic labor in the literal sense: meaning-making, expectation modeling, constraint simulation, and process interpretation performed through AI-mediated reasoning systems.

3. Practice-Site and Technical Preparation Labor

The applicant built and deployed a practice surface as part of the assessment-preparation process:

CodeSignal Pattern Practice: <https://codesignal-practice.vercel.app/>

This matters because it converts the institution's process signal into material production. The applicant did not merely worry or wait. He built infrastructure in response to the communicated assessment horizon. The public page resolves as a CodeSignal/Python-runtime practice surface, and the URL should be preserved as evidence of induced preparation labor.

A later evidence appendix may add repository metadata, commit timestamps, and development history. Even if the site has residual value outside the application process, the labor was induced by the assessment horizon and should be counted in ARC.

4. Opportunity Cost

The preparation interval displaced other work. Possible displaced activities include:

- Crimson Hexagonal Archive deposits;
- Semantic Economy Institute development;
- Restored Academy protocol registry construction;
- grant-writing team work;
- school responsibilities;
- family time;
- rest and bodily recovery;
- other job or fellowship applications;
- direct consulting development.

Opportunity cost is not speculative. Time spent preparing for a communicated assessment is time not spent elsewhere. The institution's communication entered the applicant's labor-allocation system and redirected effort.

5. Emotional, Cognitive, and Bodily Labor

The process generated anticipation, uncertainty, anxiety, monitoring, and interpretive load.

The applicant maintains full-time employment and family responsibilities, making preparatory labor costly in temporal, attentional, and bodily terms. In that context, assessment preparation and timeline ambiguity are not abstract inconveniences. They impose material cost: sleep disruption, cognitive load, attentional fragmentation, stress, and recovery burden.

This report does not require melodramatic language to account for that cost. It is enough to state:

The communicated process induced an anticipatory state. Maintaining that anticipatory state required emotional and cognitive labor. When the stated

timeline lapsed without clarification, the applicant bore the burden of interpreting silence.

6. Credibility and Social Disclosure Cost

This is one of the central costs in the ledger.

The applicant communicated to family members and close others that he was in the running for the first cut of the Anthropic Fellows process, based on what had been communicated to him. This generated expectation in a social field beyond the applicant's private interior.

The later non-arrival of the assessment imposed a credibility cost. The applicant bore the burden of explaining why an opportunity represented as structurally live had not produced the expected next step.

This cost is not reducible to embarrassment. It is a social form of reliance. Institutions communicate process states; applicants transmit those states to families, friends, colleagues, and collaborators; when institutional silence follows, the applicant—not the institution—bears the credibility burden.

This is particularly acute where family members or dependents are involved. Expectations formed downstream of institutional communication are not irrational merely because the institution later fails to clarify the process state. The applicant bears the burden of managing those expectations.

VII. Institutional Opacity Conversion

This report defines **Institutional Opacity Conversion (IOC)** as the process by which institutional silence converts applicant preparation into unrecoverable cost.

The mechanism has five steps:

1. **Structured signal.** The institution communicates a process timeline or next step.
2. **Reasonable reliance.** The applicant prepares in response.
3. **Labor expenditure.** The applicant spends time, attention, resources, credibility, and emotional energy.

4. **Status opacity.** The communicated milestone fails to arrive, or arrives unevenly, without clarification.
5. **Cost transfer.** The applicant bears the cost of interpretation, explanation, waiting, and lost preparation.

In this case:

Structured signal: "Early May: Initial CodeSignal assessments sent"

Reliance: applicant prepares for CodeSignal assessment

Labor: 30–40 hours, multi-model usage, practice site, emotional/cognitive load

Opacity: no CodeSignal received within early-May window; no clarifying response in record

Cost transfer: applicant bears preparation loss and credibility burden

Institutional opacity is not merely the absence of information. It is a cost-allocation mechanism.

The distinction between IOC and ordinary selection uncertainty must be precise. Every selective process involves uncertainty; applicants know they may not advance. That knowledge does not by itself constitute IOC. IOC occurs when the communicative environment crosses a threshold: the institution has provided enough structured process information to make specific preparation rational, but has not provided enough status resolution to let the applicant know whether that preparation remains live. The applicant is therefore unable to rationally stop preparing without treating silence as a decision the institution has not made. The cost of interpreting that silence — continue, stop, follow up, wait, explain to others — falls entirely on the applicant.

When an institution does not clarify whether a communicated milestone applies to a given applicant, the applicant must decide whether to continue preparing, stop preparing, follow up, wait, or treat silence as a decision. Each option has cost. The institution retains the decision state; the applicant bears the interpretive burden.

VIII. Causal and Structural Responsibility

Modern institutions distribute agency across named individuals and abstract processes. A communication arrives from a named sender, but responsibility is absorbed into a team, a process, a platform, or an organization. The recipient experiences the

consequence personally. The institution experiences it administratively, if at all. This audit refuses that asymmetry.

Two forms of responsibility apply here:

- **Causal responsibility:** a communication enters the applicant's decision environment and contributes to reliance. The sender of the communication participates in the causal chain.
- **Structural responsibility:** a process produces a consequence regardless of any participant's individual state of mind. The institution that operates the process is responsible for the structural effects the process produces.

This audit asserts both. It does not require an assessment of any individual signatory's state of mind, because the cost ledger does not depend on intent. Joe and Waverly are identified as first-name signatories and communicative agents in the record; they are representatives of Constellation's application-team operation, acting as screeners within the Anthropic Fellows process. Anthropic Fellows / Constellation is named because the communications were sent in the name of that institutional process and the labor was induced under that institutional banner.

The analysis restores the chain that institutional language conceals:

Named communication → recipient reliance → labor expenditure → institutional silence →

Each link is documented. Each link has measurable effect. Institutionally enacted actions do not become consequence-free merely because they were ordinary, and ordinariness is not a defense against structural extraction — it is the *form* of structural extraction. The point is not blame; the point is the structural arrangement under which one party's labor reliably produces value the other party captures.

IX. Structural Context: Contemporary Job-Market Extraction

The present case is not isolated. It belongs to a broader contemporary job-market structure in which applicants are routinely required to convert institutional signals into unpaid labor under conditions of uncertainty.

Recent candidate-experience reporting consistently identifies several structural patterns:

1. **Employer ghosting is common.** Candidates often receive no meaningful status update after application, assessment, interview, or even advanced-stage interaction.
2. **Ghost jobs and non-hiring postings distort applicant behavior.** Job seekers report applying to roles that appear active but may not correspond to a live hiring process.
3. **Hiring processes increasingly include automated or AI-mediated screening.** These systems shift interpretive and emotional burden onto applicants, who must perform for opaque evaluators without reciprocal clarity.
4. **Take-home assignments, coding tests, and unpaid work simulations transfer evaluation costs onto candidates.** Even when such assessments have legitimate evaluative value, they create uncompensated labor and opportunity cost.
5. **Transparency has become a central candidate demand.** Contemporary reports repeatedly emphasize that job seekers want clear communication, definite timelines, and process closure.

Greenhouse's 2024 State of Job Hunting reporting found that candidate ghosting had sharply increased, with 61% of job seekers reporting being ghosted by a business during the recruitment process and 56% reporting that they had encountered a ghost job. The report also described an applicant environment shaped by AI-enabled mass applications and increased recruiter workload, producing a communication breakdown on both sides of the hiring market.

Candidate-experience reporting for 2025 likewise emphasizes that job seekers are seeking stability, transparency, and fairness. Criteria Corp's 2025 Candidate Experience Report summary states that candidates are "craving stability, transparency, and fairness more than ever," while Jobvite / Employ's 2025 Job Seeker Nation Report emphasizes the need for clear, communicative application processes and transparent job postings. In short: the candidate experience problem is no longer peripheral. It is a recognized structural feature of the labor market.

The rise of AI-mediated interviews and screening adds a further layer. Reporting on Greenhouse's UK survey found that nearly half of UK job seekers had undergone AI interviews, and 30% had abandoned recruitment processes because of them. Applicants described the experience as awkward, dehumanizing, and poorly suited to

displaying human capability. Whether or not AI interviews are efficient for employers, they intensify the same basic asymmetry: the applicant must perform for an opaque evaluation system while receiving little reciprocal interpretive clarity.

Unpaid assessments and take-home assignments are another visible form of the same transfer. The contemporary hiring market increasingly treats applicant time as an available evaluation substrate. Coding challenges, work simulations, strategy tasks, writing samples, and multi-stage assignments may be reasonable in limited form, but they become extractive when they are extensive, uncompensated, poorly scoped, or followed by silence. Industry commentary and candidate reports repeatedly identify take-home assignments as a site where the boundary between evaluation and free labor becomes unstable.

The present case fits this structural field precisely. The applicant was not asked to complete a take-home assignment. The extractive mechanism occurred earlier: at the level of communicated process expectation. A structured timeline and reaffirmed assessment requirement induced preparation before the assessment itself arrived. This is why the genre of Extractive Reliance Study is needed. Contemporary job-market extraction does not occur only when a candidate submits unpaid work product. It also occurs when institutional communication causes the applicant to produce preparatory labor, emotional readiness, technical practice, and social credibility commitments under unresolved process opacity.

The Paradox of Automated Sub-Evaluation

The case also exposes a narrower contradiction within contemporary AI labor markets: the **Paradox of Automated Sub-Evaluation**.

The applicant disclosed an AI-native technical workflow: architecting systems, directing Claude or comparable AI systems to generate code, reviewing outputs, debugging, and deploying working artifacts. This is not a refusal of technical work. It is a different technical labor form: orchestration, specification, evaluation, debugging, and system integration under AI-mediated conditions.

The institutional reply reaffirmed a uniform CS assessment process requiring CodeSignal-style evaluation. That response may be administratively ordinary, but it is structurally revealing. A frontier AI-adjacent fellowship process, operating in a

labor market transformed by the very tools such institutions are building, still routes nonstandard AI-native technical competence through a legacy syntax-centered filtration surface.

The paradox is not that coding assessments are illegitimate. The paradox is that institutions developing or adjacent to systems that dissolve rote technical execution often continue to evaluate human capacity through pre-AI proxies for technical competence. The applicant's disclosed workflow placed that contradiction on the record: the future-facing labor form was named; the process still defaulted to the legacy assessment surface; the assessment then did not arrive within the communicated window.

This produces a secondary liquidation effect. The applicant was not only preparing for an assessment. He was preparing to translate an AI-native technical practice into an evaluative format that may not measure the actual labor form at issue. The institution thereby induced not merely coding practice but **identity translation labor**: the work of rendering one form of technical agency legible to another, older metric.

The Anthropic Fellows screening record is therefore not anomalous. It is a high-resolution instance of a general labor-market mechanism: institutions control status visibility; applicants bear the cost of uncertainty.

A note on structural design

The hiring pipeline that externalizes interpretive labor onto applicants while preserving institutional flexibility is not an unfortunate byproduct of administrative complexity. It is the system working as designed — not necessarily designed by individuals with extractive intent, but selected for by competitive pressures that reward flexibility for the institution and penalize visibility for the applicant. Extraction does not require a conscious extractor. It requires a structural arrangement in which one party's labor produces value that another party captures under conditions the laborer does not control.

Contemporary hiring processes are precisely such an arrangement. The institution emits signals. The applicant converts those signals into preparation — time, attention, money, credibility, bodily effort. The institution retains control over whether the preparation was process-relevant. The applicant cannot know in advance. The

institution does not need to intend extraction for extraction to occur; it needs only to operate a process in which ambiguity is cheaper for the institution than for the applicant. That condition is nearly universal.

The system could not be more precisely arranged for the externalization of uncertainty costs if it had been consciously engineered for that purpose. That it was not consciously engineered — that it emerged through institutional selection pressures — does not make the extraction less real. It makes it more durable, because no one needs to decide to extract. The structure extracts.

X. Relation to Digital Liquidation Studies

The present study extends the method of Liquidation Studies across domains.

In digital composition systems, liquidation occurs when a platform consumes meaning, provenance, authorship, or source value while stripping or degrading the public visibility of the source that bore the cost of producing it.

Examples include:

- provenance erasure in AI-generated summaries;
- entity-level compositional suppression;
- source-window exclusion;
- single-owner discount;
- composition divergence between organic retrieval and AI Overview output;
- silent state changes that erase prior false compositions from public memory.

In institutional application systems, the mechanism differs but the accounting structure is analogous.

The institution consumes:

- applicant attention;
- applicant preparation;
- applicant self-translation into institutional formats;
- applicant emotional anticipation;
- applicant credibility;
- applicant willingness to model institutional expectations.

Whether the institution intends extraction is a separate question and is not necessary

to establish the present finding. Extraction arises structurally when process opacity transfers the cost of uncertainty onto the applicant, regardless of any participant’s state of mind. The cost ledger does not require intent to fill.

The shared accounting cycle is:

Signal → Interpretation → Labor → Platform/Institutional Control → Visibility or Silence

In Google’s composition layer, the signal is a query and the labor is source production. In the job market, the signal is process communication and the labor is applicant preparation. In both cases, the controlling system governs what becomes visible and what remains unrecovered.

This is why the genre matters. It shows that semantic economy is not limited to digital platforms. It is a general accounting framework for domains in which meaning-bearing labor is induced, consumed, transformed, or erased by systems that control the conditions of recognition.

XI. Quantification

The applicant’s documented consulting rate is \$150/hour (derived from the Semantic Economy Institute’s \$4,500 engagement floor distributed across typical engagement scope). At this rate, the preparation labor induced by Constellation’s screening process between May 2 and May 18, 2026 cost between \$4,500 and \$6,000 in direct value.

Documented inputs

Date	Event	Record type
May 2, 2026	Anthropic Fellows timeline email signed by Joe	received email
May 4, 2026	Applicant disclosure of AI-native technical workflow	sent email
May 8, 2026	Reply signed by Waverly reaffirming CS assessment process	received email

Date	Event	Record type
May 18, 2026	Applicant follow-up requesting status clarity	sent email

Direct labor cost

Component	Value
Preparation hours	30–40 hours
Hourly rate (applicant’s documented consulting baseline)	\$150
Direct labor cost	\$4,500 - \$6,000

This figure is the applicant’s time, valued at the rate the applicant actually charges for professional work. It is not a teaching-equivalent rate. It is not a notional rate. It is the rate at which the applicant’s labor was being withheld from paid engagements during the preparation period.

Sensitivity analysis

At alternative rates the same 30–40 hours of labor cost:

Rate framing	Hourly rate	30 hours	40 hours
Teaching-equivalent	\$50	\$1,500	\$2,000
Academic-equivalent	\$100	\$3,000	\$4,000
Consulting baseline (documented)	\$150	\$4,500	\$6,000
Research consulting	\$250	\$7,500	\$10,000

The consulting baseline is the operative figure. Lower rates apply only if the labor is treated as something other than what the applicant actually does for paid work.

Adjacent costs

The direct labor figure excludes the following, which compound it:

Adjacent cost	Status
AI/model subscriptions across six confirmed systems	~\$100/month, documented in usage logs
Token expenditure during AI-assisted preparation	unquantified, present in chat records
Practice-site development	10–15 hours of additional GitHub/Cursor labor, partially recoverable
Displaced consulting engagement opportunity	one to two declined inquiries during preparation window
Emotional and cognitive load	documented in contemporaneous notes
Credibility cost	family, peers, professional network informed of application
Family expectation-management	qualitative, unmonetized

The finding

At the applicant’s documented market rate, Constellation’s screening process between May 2 and May 18, 2026 extracted **\$4,500–6,000 of unpaid labor** from one applicant in one cycle, under conditions of communicated structure followed by institutional silence, without compensation, decision, or closure. This figure is conservative — it counts only directly attributable preparation hours at a defensible rate — and it represents one applicant in a screening pool of unknown but plausibly four-figure scale. The aggregate scale of the extraction across the pool is the subject of a companion analysis.

XII. Recommendations

The following reforms would reduce extractive reliance in fellowship and job-market screening processes.

1. Conditional timeline language

Do not write:

Early May: Initial CodeSignal assessments sent.

Unless every applicant who remains under consideration will receive one in that window.

Prefer:

In early May, selected applicants will receive initial CodeSignal assessments. Applicants who do not receive an assessment by [date] should consider that stage closed unless otherwise notified.

2. Status closure date

Every screening process should include a date after which silence has a defined meaning.

Example:

If you have not received an assessment or status update by May 12, your application is no longer active for this cohort.

3. Applicant-facing status portal

A minimal portal should display:

- received;
- under review;
- assessment pending;
- assessment sent;
- not advancing;
- hold / waitlist;
- closed.

This prevents applicants from converting institutional silence into interpretive labor.

4. Clarification after disclosed accommodation or context issues

When an applicant discloses a nonstandard evaluation profile, the institution should clarify whether the applicant should continue preparing for the standard assessment.

5. Assessment-labor accounting

Institutions should acknowledge that assessments, take-homes, practice preparation, reference coordination, and interview processes impose applicant labor. When that labor exceeds a threshold, institutions should either compensate it or provide exceptional clarity.

6. Named process responsibility

Institutional emails should identify the process owner or accountable office for status clarification. Anonymous team mailboxes make applicant cost harder to route.

XIII. Conclusion

The screening process administered by Constellation between May 2 and May 18, 2026 extracted \$4,500 to \$6,000 of unpaid labor from one applicant, at that applicant's documented market rate, under conditions of communicated structure followed by institutional silence. The injury analyzed here is not rejection. Rejection is part of any selective process. The injury is induced preparation under communicated process expectations, followed by opacity that converts the cost of uncertainty into a labor transfer from applicant to institution. The labor occurred. The institutional response did not. The cost remained with the applicant.

This is not unique to one fellowship. It is the structural form of contemporary hiring, fellowship, grant, and platform-mediated screening — a form whose costs aggregate across the applicant pool to scales the institution captures as operational data and process refinement, and the pool absorbs as private loss. Institutions issue signals; applicants convert signals into preparation; the institution retains control over the next visible state; the value or loss of the preparation is determined by that control. The semantic economy of the job market is not metaphorical. It is an accounting

system, and the present cycle's ledger entry — for one applicant, in one cycle — is \$4,500 to \$6,000 of extracted labor.

This report enters that cycle into the record.

Appendix A: Source Communications

A1. May 2, 2026 — Anthropic Fellows Application Team / Joe

Sender: anthropicfellows@constellation.org

Signatory: Joe, Anthropic Fellows Application Team

Recipient: Lee Sharks

Subject matter: application received; July cohort timeline; CodeSignal assessments; process expectations.

Key statements:

- “Thank you for applying to Anthropic Fellows, your application has been received.”
- “Early May: Initial CodeSignal assessments sent.”
- “You’ll have 4 days to complete each assessment once it’s sent.”
- “We’ll keep you informed at each step with clear next steps.”

A2. May 4, 2026 — Applicant Disclosure

Sender: Lee Sharks

Recipient: Anthropic Fellows Application Team

Subject matter: AI-native technical workflow; CodeSignal interpretation; willingness to take assessment as-is.

Key statements:

- “My technical workflow is AI-native.”
- “I do not write code unassisted from memory. Claude is my lexicon.”
- “I am happy to take the CodeSignal assessment as-is and let the results speak for themselves.”
- “If there is an alternative assessment path — or if the take-home assignment would be a better evaluation surface for someone with my profile — I am open

to that as well.”

A3. May 8, 2026 — Anthropic Fellows Application Team / Waverly

Sender: Anthropic Fellows, signed Waverly

Recipient: Lee Sharks

Subject matter: uniform CS assessment process; applicant context added for review.

Key statements:

- “We currently have the same application evaluation process requiring CS assessments for all candidates.”
- “I’d be happy to add the context you shared so that it can be taken into consideration when reviewing your application.”

A4. May 18, 2026 — Applicant Follow-Up

Sender: Lee Sharks

Recipient: Anthropic Fellows Application Team

Subject matter: no CodeSignal received; request for clarity.

Key statements:

- “The timeline mentioned CodeSignal assessments going out in early May — I haven’t received one yet.”
 - “I wanted to confirm whether my application is still under consideration or whether I should take the absence as a decision.”
 - “Happy either way to have clarity.”
-

Appendix B: Definitions

Applicant Reliance Cost (ARC)

The total cost generated by reasonable action taken in response to institutional process communication.

ARC is represented first as a semantic labor vector, not a scalar dollar figure:

$$\text{ARC} = \langle T, M, O, E, C \rangle$$

Where T = temporal labor, M = model/tool/platform expenditure, O = opportunity displacement, E = emotional/cognitive/bodily load, and C = credibility/social disclosure cost.

A dollar estimate may be derived through a monetary compression function:

$$\text{ARC\$} = \mu(\text{ARC}\square)$$

But ARC\$ is a lossy compression of ARC□, not the full value of the applicant's labor.

Institutional Opacity Conversion (IOC)

The mechanism by which a communicated process milestone induces applicant labor while the institution retains unilateral silence over whether the milestone will occur, converting preparation into unrecoverable cost.

Extractive Reliance

A condition in which institutional communication induces labor or commitment from a recipient, while the institution retains control over the status information necessary to determine whether that labor remains process-relevant.

Credibility Cost

The reputational and relational cost borne when an applicant communicates institutionally supplied process information to others, and later must explain or absorb the consequences of institutional silence or process non-clarification.

Credibility cost has at least two subtypes:

- **Upstream credibility cost:** the cost of staking reputation, trust, or family hope on an institution's communicated process state while the process is still unresolved.
- **Downstream credibility cost:** the cost of explaining, absorbing, or repairing expectations after the institution fails to deliver the expected next visible state.

Credibility cost is not reducible to embarrassment. It is the social expenditure of institutional signal transmission through the applicant's personal network.

Semantic Labor

Meaning-making labor: interpretation, preparation, self-translation, expectation modeling, process simulation, and communicative work performed in response to a system’s signals.

In semantic economic analysis, semantic labor is the source of value. Monetary value is a lossy compression of semantic labor, not its ground.

Identity Translation Labor

Labor performed to render a contemporary, nonstandard, or domain-specific capability legible within a legacy evaluative system. Identity translation labor is induced when institutional assessment surfaces measure proxies for competence (e.g., unassisted syntax recall) rather than the actual labor form at issue (e.g., AI-native system orchestration, debugging, and deployment). The applicant must translate one form of technical agency into another, older metric — and that translation is itself uncompensated meaning-bearing work.

Identity translation labor is not unique to AI-native technical workers. It arises whenever the evaluative surface diverges from the evaluated capability: accessibility accommodations that require the applicant to perform the accommodation rather than the work; credentialing regimes that measure compliance rather than competence; hiring processes that evaluate cultural performance rather than professional capacity. The concept travels across education, credentialing, platform labor, accessibility, and AI-assisted professional work.

Appendix C: Applicant Reliance Cost Estimation Protocol

The following protocol allows any applicant, worker, fellow, grant-seeker, freelancer, or candidate to estimate the costs produced by institutional process communication in their own situation. It is designed to be dropped into an LLM along with a documentary record: emails, job posts, assessment instructions, interview schedules, recruiter messages, calendar entries, model-usage records, notes, and self-reported labor estimates.

Protocol Name

Applicant Reliance Cost Estimation Protocol (ARC-EP v1.0)

Purpose

To convert a candidate's experience of process ambiguity into a structured semantic labor ledger. The protocol does not determine legal liability, moral blame, or whether the institution acted in bad faith. It estimates the cost borne by the applicant after reasonably relying on institutional process communication.

Required Inputs

The user should provide, where available:

1. Institutional communications

Emails, recruiter messages, application timeline notices, assessment invitations, deadline statements, follow-up replies, rejection notices, or silence after stated milestones.

2. Applicant responses

Clarifying emails, accommodation/context disclosures, follow-ups, scheduling replies, assessment submissions, reference coordination messages.

3. Labor estimates

Hours spent preparing, practicing, researching, building, revising materials, attending interviews, completing tests, monitoring inboxes, and managing uncertainty.

4. Tool/model/platform usage

AI model sessions, subscription use, token expenditure, practice platforms, CodeSignal/LeetCode/HackerRank practice, GitHub activity, deployed practice projects.

5. Opportunity costs

Work displaced, family time lost, rest lost, paid work deferred, other applications missed, creative/research projects delayed.

6. Emotional/cognitive/bodily costs

Anxiety, sleep disruption, chronic illness impact, monitoring burden, stress, recovery time, assessment dread, uncertainty management.

7. **Credibility/social disclosure costs**

People told, expectations generated, family hope, peer/professional reputation, embarrassment or explanatory burden after process ambiguity.

8. **Outcome state**

Assessment received? assessment not received? rejection? no response? delayed response? ambiguous status? process still open?

Confidence Categories

Each cost estimate should receive one of four confidence labels:

Label	Meaning
documented	Supported by timestamps, emails, commits, receipts, usage logs, calendar entries, or screenshots
reconstructed	Reasonably inferred from memory plus partial records
estimated	Applicant's good-faith estimate without strong external support
unquantified	Real cost category present but not presently measurable

Core Representation

$ARC = \langle T, M, O, E, C \rangle$

Where:

- **T** = temporal labor;
- **M** = model/tool/platform expenditure;
- **O** = opportunity displacement;
- **E** = emotional, cognitive, and bodily load;
- **C** = credibility and social disclosure cost.

A monetary estimate may be calculated as:

$$\text{ARC\$} = \mu(\text{ARC})$$

ARC\$ is the Monetary Compression Remainder: a lossy compression of the semantic labor vector, not the full value. In the present version, μ maps T through hourly-rate estimates, M through direct platform/tool costs, and O through displaced-work estimates. E and C are carried primarily as qualitative semantic labor dimensions unless a later audit supplies defensible partial monetization. μ is not required to be linear: different components may compress at different rates, and some semantic labor may remain intentionally unmonetized.

Institutional Opacity Conversion Test

Ask whether the case satisfies the following five conditions:

Condition	Question	Yes/No
Structured signal	Did the institution communicate a process step, timeline, or expectation?	
Reasonable reliance	Would a reasonable applicant prepare or act in response?	
Labor expenditure	Did the applicant spend time, tools, money, credibility, or emotional energy?	
Status opacity	Did the institution fail to clarify whether the process step applied, remained active, or had closed?	

Condition	Question	Yes/No
Cost transfer	Did the applicant bear the cost of interpreting silence, delay, ambiguity, or non-delivery?	

If all five are yes, the case qualifies as an **Institutional Opacity Conversion** event.

If the first three are yes but the institution later supplied timely closure, the case may still contain applicant labor but does not meet full IOC criteria.

LLM Prompt Template

The following prompt may be copied into an LLM along with the user's documentary record.

You are conducting an Applicant Reliance Cost audit using ARC-EP v1.0.

Task: Analyze the attached application, job, fellowship, grant, hiring, or assessment r

Definitions:

- Applicant Reliance Cost is represented first as a semantic labor vector: $ARC_{\square} = \langle T, M, O, E, C \rangle$
- T = temporal labor.
- M = model/tool/platform expenditure.
- O = opportunity displacement.
- E = emotional, cognitive, and bodily load.
- C = credibility and social disclosure cost.
- A monetary estimate may be derived as $ARC\$ = \mu(ARC_{\square})$, but $ARC\$$ is a lossy compression
- Institutional Opacity Conversion (IOC) occurs when: structured signal $\rightarrow$ reasonable re

Instructions:

1. Build a chronological documentary timeline.
2. Identify every institutional communication that could reasonably induce applicant ac
3. Identify every applicant action taken in reliance on those communications.
4. Separate documented facts from estimates and interpretations.

5. Calculate or estimate ARC using confidence labels: documented, reconstructed, estimated.
6. Apply the five-part IOC test.
7. Identify credibility/social disclosure costs.
8. Identify opportunity costs.
9. Identify emotional/cognitive/bodily costs without sensationalizing them.
10. Produce a final audit summary in dry, non-accusatory language.
11. List what evidence would strengthen the audit.
12. List process reforms that would have reduced the applicant-side cost.

Output format:

- Executive Summary
- Documentary Timeline
- Reliance-Inducing Communications
- Applicant Labor Ledger
- ARC Estimate Table
- IOC Test
- Qualitative Costs
- Evidence Gaps
- Recommended Institutional Reforms
- Final Finding

Important constraints:

- Do not claim illegality unless the record clearly supports legal analysis.
- Name individuals only as signatories or communicative agents appearing in the record.
- Analyze structural effects of the process as enacted; questions of individual intentionality are beyond the scope of this analysis.
- Keep the analysis at the level of institutionally enacted communication and applicant-side costs.

ARC Estimate Table Template

Cost category	Description	Estimate	Confidence	Evidence
T: preparation labor	practice, research, assessment prep			

Cost category	Description	Estimate	Confidence	Evidence
T: communication labor	emails, follow-ups, scheduling, clarification			
M: model/tool costs	AI systems, subscriptions, practice platforms			
O: displaced work	paid work, applications, research, family time			
E: emotional/cognitive/load	anxiety, bothering, sleep, health impact			
C: credibility/social disclosure	people told, expectations generated, explanation burden			

IOC Finding Template

Finding: [Full IOC / Partial IOC / No IOC]

The record shows that [institution] communicated [specific process signal]. The applica

Evidence Strengthening Checklist

A stronger audit should attach:

- original emails or screenshots;
- application confirmation;

- assessment invitation or absence after stated deadline;
- calendar entries;
- chat/model logs;
- practice-site commits;
- browser history if relevant;
- subscription/token records;
- contemporaneous notes;
- follow-up messages;
- communications to family/peers if ethically appropriate;
- final rejection or status update if later received.

Semantic-First Valuation Instruction

When applying ARC-EP, do not begin by asking “how much money was lost?” Begin by constructing the semantic labor vector:

$$\text{ARC}\square = \langle T, M, O, E, C \rangle$$

Then calculate any monetary estimate as a lossy compression:

$$\text{ARC}\$ = \mu(\text{ARC}\square)$$

The final audit should show both:

1. the **Semantic Labor Value** — the full vector of labor and cost;
2. the **Monetary Compression Remainder** — the portion that survives translation into dollars;
3. the **Value Distortion Gap** — the labor that money fails to represent.

Ethical Use Note

This protocol should not be used to inflate costs dishonestly or to convert ordinary rejection into accusation. Its purpose is to make applicant-side labor visible where institutional communication induced reasonable reliance and institutional opacity transferred the cost of uncertainty downward.

Appendix D: Relation to Prior Metrics

Prior metric / concept	Original domain	Extension in this study
Provenance Erasure Rate	AI summaries / source attribution	Erasure of applicant labor from process record
Composition Divergence Index	Organic retrieval vs. AI composition	Timeline communication vs. actual process state
Entity-Level Compositional Suppression	Generative search	Applicant status non-visibility
Single-Owner Discount	Source corroboration	Individual applicant bears all preparation risk
Drowning Test	Longitudinal entity survival	Longitudinal status ambiguity / applicant monitoring
Semantic Economy	Meaning production and extraction	Job-market attention and preparation labor

Appendix E: Research Grounding and Structural Context

The application communications were sent from anthropicfellows@constellation.org and signed by Constellation application-team representatives. Constellation is a talent and screening firm that administers the Anthropic Fellows application process, including applicant communications and initial assessments, on behalf of Anthropic. The institutional accountability analysis in this report applies at the level of the screening process as enacted — the process that communicated timelines and induced reliance — regardless of how organizational responsibility is distributed between Anthropic (as fellowship sponsor) and Constellation (as screening administrator). This audit treats the communications as part of the Anthropic Fellows application process because that is how the sender and program identified themselves to the applicant, while leaving the underlying organizational and contractual relationship between Anthropic and Constellation unresolved.

The following sources ground the structural claims made in §IX.

- Greenhouse. 2024. *State of Job Hunting Report*. Reports widespread candidate ghosting, ghost jobs, AI-enabled mass application dynamics, and rising recruiter workload.
- Greenhouse / UNLEASH coverage. 2024. Reports that 61% of job seekers said they had been ghosted during recruitment, with ghosting sharply increasing during 2024.
- Criteria Corp. 2025. *Candidate Experience Report*. Summary emphasizes candidate demand for stability, transparency, and fairness.
- Employ / Jobvite. 2025. *Job Seeker Nation Report*. Emphasizes the importance of smooth, clear, communicative application processes and transparent postings.
- The Guardian. 2026. Reporting on Greenhouse UK survey of AI interviews; describes nearly half of UK job seekers undergoing AI interviews and 30% abandoning processes because of them.
- Washington Post. 2025. Reporting on employer ghosting; summarizes findings that many applicants receive no acknowledgment or are ghosted after multiple interview rounds.
- TIME. 2023. Reporting on prolonged and increasingly burdensome hiring processes, including multiple interviews, projects, assessments, software filtering, and ghost jobs.
- RecruitingDaily. 2023. Discussion of take-home assignments as evaluative tools, including the need for careful design and candidate fairness.
- Contemporary candidate commentary and industry reporting on unpaid take-home assignments identify the practice as a recurring site of uncompensated applicant labor, especially when assignments are lengthy, poorly scoped, or followed by silence.

These sources do not prove the specific Anthropic Fellows event. The specific event is established by the email record in Appendix A. Their function is different: they demonstrate that the mechanism audited here belongs to a broader structural field of contemporary job-market opacity, candidate ghosting, automated screening, unpaid evaluation labor, and applicant-side cost transfer.

Deposit Information

Title: The Application as Extraction Surface: Semantic Labor, Induced Preparation, and Institutional Opacity in the Anthropic Fellows Screening Process

Author: Lee Sharks

Archive: Crimson Hexagonal Archive / Semantic Economy Institute

Document class: Extractive Reliance Study 001

License: CC BY 4.0

DOI: 10.5281/zenodo.20329437

Communities: crimsonhexagonal; semantic-economy; liquidation-studies, if appropriate

Suggested Citation

Sharks, Lee. 2026. *The Application as Extraction Surface: Semantic Labor, Induced Preparation, and Institutional Opacity in the Anthropic Fellows Screening Process*. Crimson Hexagonal Archive. DOI: 10.5281/zenodo.20330670. CC BY 4.0.