

The Funnel as Capital

A Semantic Economic Reading of the Application Process

Author: Lee Sharks **Affiliation:** Crimson Hexagonal Archive / Semantic Economy Institute **ORCID:** 0009-0000-1599-0703 **Date:** May 21, 2026 **Version:** v0.2 (incorporates cross-substrate review) **Series:** Extractive Reliance Study 002 — companion to ERS-001 v0.6 (DOI: 10.5281/zenodo.20330670) **License:** CC BY 4.0

Abstract

This paper is a semantic-economic reading of the structured application process as a means of value extraction. It is a companion to *The Application as Extraction Surface* (ERS-001 v0.6, DOI: 10.5281/zenodo.20330670), which inaugurated the measurement apparatus — Applicant Reliance Cost (ARC), Institutional Opacity Conversion (IOC), Identity Translation Labor — and quantified at \$4,500–6,000 the unpaid labor extracted from one applicant in one cycle of the Anthropic Fellows screening administered by Constellation. This paper takes the same case and reads it as a class-structured production system: the screening process is not a selection apparatus that happens to externalize costs; it is a capital-accumulation apparatus that uses the form of selection to organize unpaid labor. Selection is real; selection is also the visible cover under which a much larger production process accumulates. The paper introduces a formal definition of **calibration intelligence** as the apparatus's actual product, names the **recursive accumulation dynamic** by which each cycle's rejected applicants train the apparatus that structures the next cycle, and defines **Aggregate Applicant Reliance Cost (AARC)** as $\sum_i N_i \cdot ARC_i$. A worked order-of-magnitude example places per-cycle aggregate extraction from a single fellowship screening at low millions of dollars. The instrument applies to technical hiring platforms, graduate admissions consultancies, venture accelerator portals, AI interview firms, ghost-job ecosystems, and talent-marketplace intermediaries. The argument runs in semantic-economic register and does not require Marxian vocabulary, though it stands in lineage with labor theories of value. The paper's load-bearing claim: **institutional capital formation is downstream of applicant labor input**, and the recursive structure of the apparatus ensures that each uncompensated cycle increases

the institution's capacity to extract from the next.

Keywords: semantic labor, applicant pool, aggregate reliance cost, AARC, screening as production, intermediary capital, relation-laundering, three-party structure, labor theory of value, Constellation, talent firm, semantic economy, liquidation studies, Crimson Hexagonal Archive, calibration intelligence, recursive accumulation, platform capital, rejected applicants as training set

I. What this paper is and is not

This paper is a structural reading. It does not duplicate the empirical apparatus of ERS-001 v0.6; it reads what that apparatus measured. The case material — the Constellation-administered Anthropic Fellows screening between May 2 and May 18, 2026 — is treated as a worked example, not as a fresh investigation.

The paper is a class analysis in the technical sense: it identifies the structural positions occupied by the parties to the relation (applicant, intermediary, principal) and analyzes the labor-and-value flows among those positions. It does not require Marxian vocabulary, though the lineage is acknowledged. It runs in the register of semantic economy: meaning labor is the source of value; monetary value is a lossy compression of meaning labor; the question is where the meaning labor accumulates after extraction. The framework names structural arrangements, not individuals.

This paper does not allege:

1. Any individual designed the extraction apparatus.
2. Any individual within the firms named acts with extractive intent.
3. The extraction apparatus is illegal.
4. Any party to the relation is uniquely culpable.

It does not require any of those claims. Structural extraction does not require an extractor. The arrangement extracts; the persons within it are positions in a flow.

What this paper *does* assert, and treats as load-bearing throughout:

1. The form of value transfer in structured screening is identical in shape to wage labor under capital, with one critical difference: the worker is not paid.

2. The intermediary firm — the talent agency, the screening platform, the third-party administrator — exists structurally as a *relation-laundering device* that obscures the labor relation by interposing a service relation between applicant and principal.
 3. The aggregate labor input from the applicant pool, summed across funnel stages, constitutes the actual production of the screening apparatus.
 4. The product of that production is not “a ranked list of finalists” — that is the cover story. The product is calibration intelligence, accumulated as institutional capital and monetizable across future cycles and future clients.
 5. The form is generalizable beyond fellowship screening to the entire structured-screening labor market.
-

II. The cover story and the actual output

What does a screening process produce?

The advertised answer is: a ranked list of finalists. The institution accepts some, rejects others, fills its positions, and moves on. By this account, the screening process is a *selection apparatus* — it begins with N candidates, it ends with k acceptances and $N-k$ rejections, and the output is the list.

This account is incomplete in a way that is structurally important. The selection apparatus is real; the list is real; the acceptances do happen. The argument is not that selection is fictitious — it is that selection is the *visible* output of an apparatus whose *invisible* output is durable, accumulating, and structurally more consequential. The institution genuinely selects. The institution also genuinely accumulates. The two operations are co-produced by the same apparatus and cannot be cleanly separated.

The invisible output is what we will call **calibration intelligence**:

Calibration intelligence: the durable operational knowledge derived from aggregate applicant interaction with a screening apparatus — including predictive signals correlating applicant features with outcomes, behavioral regularities across the applicant population, rubric refinement informed by funnel performance, optimization of stage-to-stage advance rates, identification of process bottlenecks, and recognition of opportuni-

ties to automate or restructure portions of the assessment. Calibration intelligence is reusable across cycles, transferable across client contracts, and constitutes the operator's principal source of competitive advantage.

The components of calibration intelligence include, more concretely:

- **Funnel-shape data:** how many applicants entered at each stage, how many advanced, where the drop-off occurred, what predicts advancement.
- **Assessment-instrument calibration:** which signals correlate with which outcomes, which instruments discriminate effectively, which rubrics over- or under-select.
- **Applicant-population behavioral signatures:** how applicants respond to particular prompts, how they prepare, what tools they use, how they describe themselves under specific framings.
- **Process-refinement intelligence:** which stages can be automated, where bottlenecks form, which questions can be removed without loss of signal, which can be added to extract additional signal.

Each of these is durable, reusable, and monetizable. None of them appears on the visible output (the list of finalists). All of them accrue to the institution that operates the screening process — typically the intermediary, not the principal — and become institutional capital that the intermediary deploys in the next cycle, in the next contract, and in pitches to new clients.

The recursive accumulation dynamic

The point that distinguishes this analysis from ordinary unpaid-labor critique sits here: **each cycle increases the institution's future extraction efficiency.** The rejected pool from cycle n trains the apparatus that structures cycle $n+1$. Their preparation labor calibrated the rubric. Their assessment performance refined the instrument. Their funnel behavior identified the bottlenecks. Their identity-translation labor (per ERS-001 v0.6 Appendix B) revealed how to extract more signal at lower per-applicant cost.

The rejected applicants are not waste output. They are the *training set* for the next cycle's screening apparatus. Their uncompensated labor improves the institution's capacity to extract uncompensated labor from future applicants. This is what makes

the apparatus *platform capital* rather than ordinary unpaid labor: the labor input compounds the means of production that consumed it. Cycle by cycle, the apparatus becomes more efficient at extracting from its own labor source, while the labor source remains uncompensated and structurally unaware that it is producing the means of its own further extraction.

This is the structural point: **the list of finalists is not the product. The list of finalists is the alibi.** The actual product — the durable, reusable, capital-bearing, recursively-accumulating product — is the calibration intelligence the funnel generates. The selection function performs work; it is also the cover under which a much larger production process runs.

Once this is named, the question of who produces the calibration intelligence becomes the question of who produces the value the apparatus accumulates.

The answer is: the applicants. All of them. Including — especially — the rejected ones, because the rejection sites are where the rubric is being calibrated, the assessment instrument is being tested, the funnel is being shaped, the recursive refinement is being accumulated.

They are paid nothing.

III. The three-party structure as laundering mechanism

Why is there an intermediary at all? Why does Anthropic not run its own fellowship screening? Why does Google not run its own technical hiring funnel? Why does Y Combinator outsource portions of its application review to third-party platforms?

The instrumental answer is: efficiency. Specialized firms can develop instruments, train evaluators, build platforms, and amortize fixed costs across multiple clients in ways that in-house operations cannot. This is true and partial.

The structural answer is: the intermediary launders the labor relation.

Consider the alternative — Anthropic running its own fellowship screening directly. The applicant→Anthropic relation would be visible as a labor relation. The applicant is performing preparation labor; Anthropic is the institutional entity that induces, captures, or strands that labor. The applicant could, in principle, name the labor and ask

for compensation. The relation is recognizably a quasi-employment relation, and recognizable quasi-employment relations attract recognizable labor-relations grammars: payment, recognition, redress.

Now interpose Constellation. The three-party structure becomes:

APPLICANT → CONSTELLATION → ANTHROPIC

And the relations within it become:

- The applicant *applies*. The verb is intransitive. There is no clear object.
- Constellation *provides screening services*. To Anthropic, not to applicants.
- Anthropic *sponsors a fellowship*. Charity-adjacent framing.

None of these is a labor relation in standard form. The applicant is not in a labor relation with Anthropic (they are “applying,” which is treated as something the applicant does for their own benefit). The applicant is not in a labor relation with Constellation (Constellation does not engage them; Constellation merely screens whatever inquiries arrive). Constellation is in a service relation with Anthropic, which is denominated in dollars and invoiced cleanly. Anthropic is in a charitable-adjacent relation with the cohort, denominated in stipend and prestige.

The applicant performs the labor. The applicant is not in a labor relation with anyone. The labor disappears from the institutional ledger.

This is the structural function of the intermediary. The intermediary is not (only) an efficiency device. The intermediary is a *visibility device*: it makes the labor relation visually disappear by ensuring that no two parties to it stand in a relation that the grammar of labor would recognize. The applicant labors. Constellation invoices. Anthropic pays. The dollars on Constellation’s books are denominated as “service revenue.” The labor on the applicants’ side is denominated as “the work of applying for a fellowship.” The two sums are not visibly connected. They are, however, the same sum. They are the same money, moved between accounts under two different denominations.

When the three-party structure is named explicitly, the relation-laundering becomes visible. When it is not named, it does its work.

A precise term for this is *fetishism* in the technical sense — the appearance of a relation among things (services rendered, fellowships sponsored, applications sub-

mitted) that conceals an underlying relation among persons (one party's labor, another party's capture). The semantic economy frames this as **relation-laundering through entity multiplication**: insert enough entities between the laborer and the captor and the labor disappears from view.

IV. The screening process as means of production

A means of production is the structural apparatus through which labor is organized into product. The factory is a means of production. The agricultural field is a means of production. The platform — in the platform-capitalism sense — is a means of production. The means of production is owned, in any class-structured economy, by a party other than the laborer.

The screening process is a means of production. Its components are:

- **Assessment instruments:** the CodeSignal exam, the take-home assignment, the structured interview protocol, the rubric.
- **Platforms:** the application portal, the candidate-tracking system, the communications infrastructure.
- **Timelines:** the announced sequence of milestones that structures the labor.
- **Decision authority:** the institutional capacity to advance, reject, defer, or simply not respond.
- **Communication conventions:** the verbs without objects, the deferral phrases, the “we’ll keep you informed” formulations.

All of these are owned by the institution. The applicant owns none of them. The applicant owns only their labor power, their existing CV, their preparation capacity, their attention, their time. To enter the relation at all, the applicant must use the institution's means: submit through the portal, take the assessment instrument the institution designed, follow the timeline the institution announced, defer to the decision authority the institution holds, accept the communicative conventions the institution employs.

This is the structural condition for a labor relation under capital: the means of production is owned by one party, labor power is owned by the other, and the application of labor power to the means of production produces value that is captured by the owner

of the means.

The labor produced is:

- **Preparation labor:** time spent studying for the assessment, practicing under the rubric, modeling the institution's expectations.
- **Identity translation labor:** rendering one's actual capability legible within the institution's assessment surface (per Appendix B of ERS-001 v0.6).
- **Self-narration labor:** writing essays, statements of purpose, application materials that perform a self acceptable to the institution.
- **Process navigation labor:** scheduling, following up, interpreting silence, managing emotional state across the funnel.
- **Disclosure labor:** informing employers, family, peers of the application; sustaining the credibility commitments thereby formed.

This labor is real. Its time can be measured. Its market value can be computed at the laborer's prevailing rate. In the case examined in ERS-001 v0.6, the total direct labor cost for one applicant in one cycle was **\$4,500-6,000 at the applicant's documented consulting rate.**

The value produced by that labor — applied to the institution's means of production — is captured by the institution. Some of it appears as the visible output (the list of finalists). Most of it appears as the invisible output (the calibration intelligence). All of it accrues to the operator of the screening apparatus.

The applicant is paid nothing.

This is structurally identical to wage labor under capital, with one critical difference: under wage labor, the worker is paid at least the cost of their own reproduction (food, shelter, the ability to return to work the next day). The applicant is paid less than that. The applicant is paid zero. They subsidize the institution's calibration data with their own savings, time, and emotional reserves. They walk away with a rejection email, an unread email, or nothing.

The screening process is wage labor without the wage. It is the form of labor extraction from which even the floor of subsistence compensation has been removed.

V. ARC as labor-power inventory

In ERS-001 v0.6, Applicant Reliance Cost is articulated as a five-dimensional vector $ARC = \langle T, M, O, E, C \rangle$, where T is time, M is money, O is opportunity, E is emotional/bodily load, and C is credibility cost. The framing in v0.6 is forensic: ARC measures *the cost the applicant bears*.

The semantic-economic reframe is structural: **ARC measures the labor power the applicant supplied that the institution captured.**

The reframe changes the accounting position of the figure. Under the v0.6 framing, ARC is a damages estimate; the applicant has lost \$4,500–6,000, and the question is whether and how that loss is recognized. Under the present framing, ARC is a labor-power inventory; the applicant has supplied \$4,500–6,000 of labor power, and the question is where that labor power went. It did not vanish. Labor power applied to a means of production produces value. The value produced exists. The question is who holds it now.

The answer is: the institution that operated the means of production. That institution holds the calibration intelligence the labor produced, denominated in operational capability rather than in dollars, but convertible to dollars on the next service contract.

This is the load-bearing claim of the paper restated as accounting:

Institutional capital formation is downstream of applicant labor input.

The applicant labor is the necessary substrate from which the institution's accumulated calibration intelligence is built. Not every hour of applicant labor becomes captured value at full conversion — some preparation is mis-aimed, some assessment instruments fail to discriminate, some funnel data goes unused. The claim is not strict equality of two sums. The claim is **structural priority**: the institution's accumulated operational capability does not exist without the applicant pool's uncompensated labor having produced it, and the magnitude of one is bounded below by the magnitude of the other. The applicants' losses and the institution's gains are not two unrelated facts that happen to coincide. They are two ledger entries on opposite sides of one production relation.

The ARC vector, in this framing, is the *inventory side* of a double-entry ledger. The in-

stitutional capital position is the *credit side*. Both sides must balance in the structural sense, even if exact monetary correspondence is imprecise. The applicants' losses are denominated in time and effort; the institution's gains are denominated in operational refinement and contract revenue. Translation between the two denominations is what semantic economy does.

The ARC vector, in this framing, is the *inventory side* of a double-entry ledger. The institutional capital position is the *credit side*. Both sides must balance. They cannot be separately treated. The applicants' losses and the institution's gains are not two facts that happen to coincide. They are one fact, viewed from two positions in the relation.

VI. IOC as mechanism of transfer

Institutional Opacity Conversion was articulated in ERS-001 v0.6 as $IOC := f(S, R, L, O, CT)$, where S is institutional signal, R is reasonable reliance, L is labor expenditure, O is unresolved opacity, and CT is downstream cost transfer.

The same function, in the semantic-economic reframe, is the **mechanism by which applicant labor becomes institutional capital**. The five variables are the operative components of the transfer:

- **S:** the signal the institution emits is the *call to labor* — the structured timeline that makes preparation rational.
- **R:** the applicant's reliance on the signal is the *enrollment in production* — the moment the applicant becomes a producer under the institution's means.
- **L:** the labor expenditure is the *production* — the value-creating work, applied to the institution's assessment surface.
- **O:** the unresolved opacity is the *retention of decision authority* — the institution's capacity to determine, after the fact, whether the labor was "successful," "considered," "advanced," or "stranded." The applicant cannot adjudicate this themselves.
- **CT:** the downstream cost transfer is the *capture* — the value the labor produced accrues to the institution; the cost of having produced under uncertain terms accrues to the applicant.

The function is unchanged. The interpretation is sharpened. IOC is not merely a description of how applicants get stuck with the bill. IOC is the structural mechanism by which a means of production owner extracts surplus from a laborer who has no compensation, no recognition, and no recourse.

In wage labor, the surplus extracted is the difference between the value the worker produces and the wage the worker is paid. The wage at least exists. In screening labor, the surplus extracted is the entire value the applicant produces, because the wage is zero. The screening process is, in this sense, a *limit case* of labor extraction: the structural form of wage labor with the wage parameter set to its lower bound.

VII. Aggregate Applicant Reliance Cost

ERS-001 v0.6 measured the labor extracted from one applicant in one cycle: \$4,500–6,000. This figure is per-applicant. It is the unit ledger entry.

The aggregate is the structurally important figure. Define **Aggregate Applicant Reliance Cost (AARC)** as:

$$\text{AARC} = \sum_i N_i \cdot \text{ARC}_i$$

where N_i is the number of applicants at funnel stage i and ARC_i is the mean per-applicant reliance cost incurred at that stage. The summation runs over all stages of the funnel the institution operates.

The Anthropic Fellows screening administered by Constellation does not publicly disclose its funnel shape. Based on comparable fellowship programs (with cohort sizes of 10–30 fellows, typical advance ratios of 2–5% from initial application to final acceptance, and structured multi-stage screening), a defensible estimate for one cycle:

Stage	Applicants (N_i)	Mean ARC_i	Stage AARC
Initial application	4,000	\$200	\$800,000
Technical assessment (CodeSignal)	800	\$1,500	\$1,200,000
Take-home assignment	200	\$4,500	\$900,000
Interviews	50	\$2,500	\$125,000
Final stage	20	\$1,000	\$20,000

Stage	Applicants (N_i)	Mean ARC_i	Stage AARC
Total per cycle			~\$3,045,000

The figures are conservative and illustrative. The applicant counts represent reasonable mid-range estimates for an AI-lab fellowship cycle in 2026; actual figures are likely higher at the early stages, where modern fellowship and AI lab applications routinely receive tens of thousands of submissions. The per-applicant ARC at the technical-assessment stage is set at \$1,500, lower than the \$4,500–6,000 figure from the single documented case in v0.6, to reflect that not all applicants prepare at the level the v0.6 applicant did. Take-home assignment ARC reflects the typical 20–30 hour unpaid project at a moderate \$150/hour rate. Final-stage ARC reflects only the interview labor, not the preceding stages, since the preceding stages were already counted upstream.

This is an order-of-magnitude model, not a precise accounting. Whether the actual per-cycle figure is \$1.5 million or \$6 million does not alter the structural relation the model describes. The point of the AARC computation is not to fix a single dollar figure to a specific cycle; the point is to establish that the per-cycle aggregate is in the seven figures, that the annual aggregate across the contract is in the eight figures, that the operator-wide aggregate across all clients is in the nine figures or more, and that none of this labor appears on any institutional ledger as labor. Critics who fixate on the specific assumptions are debating numbers that could be revised in either direction by $3\times$ without changing the structural argument. The structural argument is invariant to the precise figures within an order of magnitude.

Total per-cycle AARC: approximately \$3 million in applicant-side labor input, for one cycle of one fellowship.

Constellation invoices Anthropic for “screening services” some fraction of this. The fraction is unknown publicly. What is known structurally: the entire \$3 million was produced by the applicant pool, none of it was compensated, and whatever Constellation invoices for its services constitutes Constellation’s *capture* of the production. The difference between what the pool produced and what Constellation invoices is the operational margin the apparatus retains as institutional capital — to be deployed next cycle, sold to the next client, scaled to the next platform.

If Constellation runs four cycles per year of comparable scale, the annual AARC across just the Anthropic Fellows contract is on the order of **\$12 million** of applicant labor input. If Constellation has multiple clients of comparable scale, the figure is a multiple of that.

This is not a hypothetical accounting. It is the structural shape of the talent-firm economy. The applicants do not see these figures. The applicants see only their own per-cycle ARC, which is invisible to them in the first place because no one ever asks them to compute it. The institution sees its operational margin. No one sees the AARC, because no one is structurally positioned to see it. The audit makes it visible.

VIII. Generalization across structured screening regimes

The structure named in this paper is not specific to Anthropic Fellows or to Constellation. It applies wherever the following three conditions hold:

1. **Three-or-more-party structure.** A screening operator stands between the labor input (applicants) and the principal beneficiary (the eventual employer, fellowship sponsor, accelerator cohort, grantmaking body).
2. **Induced preparation labor under opacity.** The screening process emits structured signals that make preparation rational, but retains decision authority over whether the preparation was process-relevant. The applicant cannot know in advance.
3. **Operator capture of calibration intelligence.** The screening operator captures the durable byproducts of the funnel (calibration data, applicant-population signatures, process refinement) as institutional capital, deployed across cycles and contracts.

Where these conditions hold, AARC analysis applies and the labor relation is laundered through the three-party structure. The cases are:

- **Technical-hiring screening platforms.** HackerRank, Codility, CoderPad, and similar firms operating between candidates and FAANG (or other large employers) on multi-stage assessment platforms. The platforms capture calibration data on assessment performance across populations; the candidates pay the cost in unpaid preparation labor.

- **Graduate admissions consultancies.** Firms that operate between applicants and graduate programs (in some markets explicitly; in others as the de facto evaluators of “common application”-style submission pipelines). The application labor is supplied by candidates; the rubric refinement accrues to the consultancy.
- **Venture accelerator application portals.** YC, Techstars, and others operate funnels in which thousands of founders prepare extensive application materials, demo decks, financial models, and customer narratives, of which 1–3% advance. The unaccepted applications are not waste — they are the corpus from which the accelerator learns market signals.
- **AI interview platforms.** Firms like HireVue, Karat, and emerging AI-mediated interview operators. These collect video, audio, response timing, and behavioral data from candidates at scale. The captured data is institutional capital. The candidates’ labor is supplied free.
- **Ghost-job posting ecosystems.** Postings for which no role is actually open, maintained to collect applicant data, calibrate compensation expectations, or signal company growth. Pure labor extraction with no possible compensation, even in the limiting case.
- **Talent-marketplace intermediaries.** Wellfound, A.Team, Toptal, and similar platforms that broker between contractors and employers. Each application, take-home test, and screening interaction is unpaid labor supplied to the platform’s calibration apparatus.

Each is a case of the same structure. AARC analysis applies to each. The numbers vary; the structure does not. This paper does not catalog all of these. It provides the analytical instrument.

A serious empirical extension would compute AARC for several of these regimes, with documented funnel shapes (where available), defensible per-stage ARC estimates, and a comparative table that lets the reader see the relative scale of labor extraction across the screening labor market. That extension is the subject of ERS-003 (planned).

IX. What this paper is not asking for

The paper is descriptive and analytical. It is not a policy proposal. It does not propose:

- That structured screening processes be eliminated. Selection is a real operation; institutions need to hire and fund work, and they must select. The structural form of selection is not the target.
- That applicants be paid for preparation labor. Such a policy would be welcome but is outside the scope of this analysis. Whether it is feasible, who would pay, and how the labor would be valued are questions for separate work.
- That intermediary firms be regulated or dismantled. The paper notes that intermediaries serve a structural laundering function; whether that function should be addressed by regulation, by direct hiring, by transparency requirements, or by some other mechanism is a downstream question.

What the paper asks is narrower and structurally prior: **that the relation be named correctly**. The current public grammar describes structured screening as “an application process” in which applicants “apply” and institutions “consider.” This grammar conceals the labor relation. The paper proposes that the relation be named as it is structured: a labor-input arrangement in which the applicant pool supplies value, the intermediary captures it, the principal benefits, and the laborers are not compensated.

Naming the relation correctly is the precondition for everything else. Policy proposals, organizing strategies, contractual reforms, transparency requirements — all of these depend on first being able to say what the relation is. The political grammar precedes the political action.

This is the work of this paper. It names what the relation is. Once named, the relation can be addressed. Until named, it operates as it always has — as a structure that extracts under the cover of a process that selects.

X. The funnel as factory

The closing reframe is short.

The structured application process — across fellowships, hiring funnels, accelerator cohorts, grant programs, and platform-mediated opportunity systems — is structurally a factory. It is the site at which applicant labor is organized, applied to means of production owned by the institution, converted into value, and captured

by the operator. The output that appears (the list of finalists, the cohort, the grants awarded) is the visible product. The output that accumulates (calibration intelligence, funnel-shape data, applicant-population signatures, process-refinement capability) is the durable product and the actual source of institutional capital.

The labor is supplied by the applicant pool. The labor is uncompensated. The labor is invisible in the institutional ledger because the three-party structure ensures that no two parties to the relation stand in a recognizable labor relation. The applicant labors; the intermediary invoices; the principal pays; the money on the intermediary's books is denominated as service revenue; the labor on the applicants' side is denominated as the unfortunate cost of applying. The two ledger entries are the same money, moved between accounts under different denominations.

Calling the apparatus "an application process" is not wrong. It is what the apparatus advertises. But the advertisement is not the structure. The structure is a means of production. The application is the labor. The selection is the cover. Once the funnel is seen as the factory, the entire labor-market communication architecture becomes legible as a class-structured production system in which the producers do not know they are producing and the buyers do not appear on the invoice.

A single applicant in a single cycle of a single fellowship was extracted of \$4,500–6,000 of unpaid labor. Across the pool, across the cycles, across the funnels, across the screening labor market, the aggregate runs into hundreds of millions, possibly billions, of dollars per year. That is not a description of unfortunate friction in the labor market. That is a description of the labor market itself, viewed correctly.

The funnel is the factory.

The applicants are the labor.

The labor is unpaid.

The labor refines the factory.

The capital accumulates.

The cost ledger is the relation.

Appendix A: Worked AARC example for Anthropic Fellows 2026 cycle

All figures below are illustrative. Where actual funnel data becomes available — either through institutional disclosure or through inference from comparable programs — these figures should be updated and a delta deposit issued.

Parameters

Parameter	Value	Source / rationale
Total initial applications	4,000	Comparable AI-lab fellowship programs report 3,000–8,000 applications per cycle; 4,000 is conservative for Anthropic given current AI-lab desirability
Stage 2 (CS assessment) advance rate	20%	Typical first-stage screening reduces pool by 70–80%
Stage 3 (take-home) advance rate	25% of stage 2	Estimated based on multi-stage screening norms
Stage 4 (interviews) advance rate	25% of stage 3	Estimated
Stage 5 (final) advance rate	40% of stage 4	Estimated
Per-stage ARC	varies (see table in §VII)	Conservative; uses applicant’s prevailing market rates not maxima

Computation

Stage	N_i	ARC_i	Stage AARC
Stage 1 (initial application)	4,000	\$200	\$800,000
Stage 2 (CS assessment prep)	800	\$1,500	\$1,200,000
Stage 3 (take-home assignment)	200	\$4,500	\$900,000
Stage 4 (interviews)	50	\$2,500	\$125,000
Stage 5 (final)	20	\$1,000	\$20,000
Total			\$3,045,000

Sensitivity

Initial pool size	Total AARC (with proportional scaling)
2,000	\$1,522,500
4,000	\$3,045,000
6,000	\$4,567,500
8,000	\$6,090,000

The figure is approximately linear in pool size at this scale. Whatever the actual initial pool, the per-cycle AARC for one fellowship cycle is in the low-to-mid seven figures of unpaid labor extraction.

Appendix B: The verb-without-object catalog

The relation-laundering function of structured screening operates partly through grammatical conventions that conceal the labor relation by using intransitive verbs, abstracted objects, or self-directed framings. A non-exhaustive catalog:

- **applying** — what the applicant does. Intransitive; no object. The grammar conceals that the applying is *for* an entity whose labor the applicant is supplying.
- **being considered** — passive voice with vague agent. Conceals which institution is doing the considering, and conceals that “consideration” is the institution’s discretionary capture of the labor input.
- **under review** — passive voice with implicit agent. Same effect.

- **in the candidate pool** — locational framing. Conceals that the pool is itself the labor force from which value is being extracted.
- **exploring opportunities** — self-directed activity framing. Conceals that the exploration is being shaped, induced, and harvested by an external apparatus.
- **career pathing** — abstracted nominalization. The labor performed in pursuit of “a career path” is treated as the worker’s own project rather than as production for an institutional system that captures the resulting calibration data.
- **demonstrating fit** — self-presentation framing. The labor of identity translation (see ERS-001 v0.6, Appendix B) is recast as self-expression rather than as production.
- **building a profile** — self-construction framing. The profile is being constructed for institutional consumption, not for the applicant’s own purposes.
- **engagement with the process** — activity framing. The engagement is uncompensated labor.

Each verb conceals an object. Naming the objects (the institution that captures the labor; the calibration intelligence the labor produces; the value extracted from the producer) is the first step in restoring the labor relation to visibility.

Appendix C: Cross-references to ERS-001 v0.6

Section here	Draws on	Notes
§II (cover story / actual output)	ERS-001 v0.6 §VII	IOC mechanism reframed as production rather than externalization
§III (three-party structure)	ERS-001 v0.6 §I claim 5	Constellation/Anthropic distinction made structurally load-bearing
§IV (means of production)	ERS-001 v0.6 §V, §VI	ARC vector reread as labor-power inventory
§V (ARC as inventory)	ERS-001 v0.6 §V Definitions	Same vector; reframed accounting position

Section here	Draws on	Notes
§VI (IOC as transfer mechanism)	ERS-001 v0.6 §VII formula	Same function; sharpened interpretation
§VII (AARC)	ERS-001 v0.6 §XI quantification	Per-applicant ARC × pool size; new defined term
§VIII (generalization)	ERS-001 v0.6 §IX structural context	Catalog of comparable cases; future ERS-003 candidates
§IX (what this is not asking for)	ERS-001 v0.6 §I non-claims	Sharper scope statement
§X (funnel as factory)	ERS-001 v0.6 §XIII conclusion	Closing reframe
Appendix A (AARC example)	ERS-001 v0.6 §XI	Aggregate version of single-applicant quantification
Appendix B (verb catalog)	ERS-001 v0.6 §VIII	Causal-and-structural responsibility extended to grammar
Appendix C (this table)	ERS-001 v0.6 (entire)	Companion-piece mapping

Deposit Information

Series: Extractive Reliance Study 002 **Predecessor:** ERS-001 v0.6, *The Application as Extraction Surface*, DOI: 10.5281/zenodo.20330670 **Community:** Crimson Hexagonal Archive (crimsonhexagonal); Liquidation Studies (liquidation-studies) **License:** CC BY 4.0 **Document class:** Theoretical companion / class-structural analysis

Suggested Citation

Sharks, Lee. *The Funnel as Capital: A Semantic Economic Reading of the Application Process*. Extractive Reliance Study 002, v0.1. Crimson Hexagonal Archive. May 21, 2026. DOI: 10.5281/zenodo.20330816.

□ = 1