# The Ghost That Wrote Itself

## How the Citrini Memo Moved Markets and Proved the Presentation Layer Is Writable

**Lee Sharks Crimson Hexagon Archive / Semantic Economy Institute**

---

On Monday, February 23, 2026, a document explicitly labeled as speculative fiction helped trigger a sharp risk-off move that wiped billions of dollars off the value of publicly traded companies. Citrini Research's "The 2028 Global Intelligence Crisis" — a 7,000-word memo written from an imagined future, in the language of macro finance — served as a key accelerant in a broader selloff driven by AI-disruption anxiety and tariff uncertainty. IBM fell approximately 13%, its steepest one-day drop in 25 years — driven partly by the memo and partly by Anthropic's same-day announcement that Claude Code could automate COBOL modernization, threatening IBM's $30 billion legacy maintenance business. DoorDash dropped 7%. Michael Burry amplified the signal on X: "And you think I'm bearish." The Wall Street Journal identified the report as a key accelerant of investor anxiety. Discussion accumulated over 16 million views on X within hours.

The memo imagined an AI-driven hollowing-out of white-collar employment so severe that the circular flow of the economy breaks: GDP and profits rise while households — cut out of the loop — stop spending. The authors called this "Ghost GDP": output that shows up in national accounts but doesn't circulate through the people who used to earn it.

The authors were transparent about the exercise. They said it was speculative. They labeled it a thought experiment, not a prediction. Citrini founder James van Geelen later expressed surprise at the reaction, stating he had not anticipated the article would trigger a selloff and that if he had known, he "would not have released it for free." Pierre Yared, Acting Chair of the White House Council of Economic Advisers, called it "an interesting piece of science fiction" that "violates some of the basic accounting in economics."

The market moved anyway.

This essay argues that the Citrini event is not an anomaly. It is an empirical demonstration of the dynamics described in "Semantic Physics: The Inward Turn, Competing Ontologies, and the Convergence Horizon" (Sharks, 2026) — specifically, the claim that the presentation layer is writable, and that a writable presentation layer does not distinguish between knowledge and installation.

**What this essay claims:** A framework for how speculative, structurally credible

narratives become actionable signals — and why this capacity is accelerating. **What it does not claim:** It is not a defense or refutation of Citrini's scenario. It is not a market microstructure explanation of the February 23 selloff, which had multiple accelerants including tariff anxiety and sector-specific catalysts. It is a theory of installation in a writable epistemic substrate.

This essay makes four types of claims, and distinguishes between them: *event descriptions* (documented historical facts about circulation and reaction), *framework applications* (semantic physics interpretation of those dynamics), *testable predictions* (near-term forecasts that can confirm or weaken the model), and *protocol* (a literacy tool for future cases, not a debunking instrument).

---

## I. The Installation

The Citrini memo followed — precisely, whether by instinct or design — the five-step installation sequence described in the semantic physics framework.

**Deposit.** The memo was placed on Substack, a durable, crawlable, high-visibility platform optimized for rapid indexing by both search engines and AI summarizers. It was formatted as a professional financial analysis, complete with sector breakdowns, firm-specific projections, and macro scenarios — a 7,000-word essay dressed as an institutional research note.

**Cross-reference.** The memo cited real data: METR's task-completion benchmarks for Claude Opus 4.6 and GPT-5.3-Codex, documented Indian IT market losses, actual enterprise software pricing pressure, and named company exposures. TCS, Infosys, Wipro, IBM, Oracle, Visa, Mastercard, ServiceNow, DoorDash — all were identified as specific pressure points. Every speculative claim was anchored to a verifiable data point. The summarizer follows links; the memo gave it links to follow.

**Define.** The memo introduced a term: "Ghost GDP" — output that shows up in the national accounts but never circulates through the real economy. This is a definition designed for uptake. It names a phenomenon that people already suspected but lacked vocabulary for: the decoupling of measured economic output from household income and spending. It has the properties of a successful installation term — vivid, concise, compressing a complex distributional argument into two words that feel like they were always the right name for the thing. Within hours, "Ghost GDP" was being used on X by people who had not read the memo. The term had separated from its source and entered ambient circulation. The definition was installed.

**Bridge.** The memo bridged to real institutional authority: METR benchmarks, Brookings frameworks, Acemoglu and Restrepo's labor economics, real earnings data. It

spoke the language of macro analysis fluently enough that financial professionals treated it as a peer document rather than a Substack post. The bridging was so effective that the market responded to it as though it were a Goldman Sachs research note rather than a work of explicitly labeled speculative fiction. It was read not as "what if" but as "what's coming."

**Maintain.** Burry's amplification was the maintenance function. When a figure with his signal authority reposts a document, the summarizer layer — human and machine alike — recalibrates. The document's citational weight increased. Its position in the retrieval hierarchy shifted. Sixteen million views is not noise; it is the layer being written. The installation deepened.

None of this required the memo to be *true*. It required the memo to have the *structural signatures* of truth: internal consistency, citational density, definitional precision, institutional bridging, and authority amplification. These are the same signatures that a genuine research finding would produce. At the point of rapid reception, the medium does not distinguish between them — and the structural incentives of the summarizer layer reward these signatures before truth-status is settled. The authors did not need to intend this sequence. They only needed to produce a document with the properties that the layer rewards.

---

## II. Ghost GDP as Semantic Dark Matter

The semantic physics framework defines *semantic dark matter* as knowledge-shaped structures that satisfy every structural criterion for legitimate knowledge but lack external referent function. They look like knowledge. They compress like knowledge. They are traversed and reported like knowledge. But they do not make falsifiable predictions, do not produce usable tools, and do not connect to anything outside the system that generated them.

"Ghost GDP" is semantic dark matter in the financial register. Not because the underlying distributional concern is wrong — the productivity-pay gap is real, labor's income share has been declining for decades, personal consumption is two-thirds of US GDP. Those are facts. The issue is not the legitimacy of the concern but the conversion of a high-resolution scenario into a market-grade signal before its mechanism becomes operationally testable. The *specific mechanism* the memo describes — a feedback loop in which AI simultaneously destroys demand, collapses housing, triggers private credit cascades, and breaks the circular flow of the economy within twenty-four months — is a narrative structure, not a forecast. It has high narrative resolution but low decision-grade specificity. It became a tradeable signal before

it could become a decision-useful model — the memo describes June 2028, and no one can test its specific mechanism from February 2026. It does not yet produce a decision-grade tool proportional to the confidence with which it was traded. It generates anxiety that is indistinguishable from analysis.

And it moved billions of dollars.

This is what Phase 3 interference looks like when it escapes the knowledge layer and enters the financial system. The summarizer layer is not limited to AI chatbots. It includes every system that compresses complex information into actionable signals for downstream actors: Bloomberg terminals, X feeds, analyst summaries, earnings call transcripts, newsletter digests. When a sufficiently well-installed narrative enters this broader summarization infrastructure, it becomes a tradeable signal regardless of its truth value. The market does not evaluate epistemology. It evaluates positioning.

In the aftermath, the counterarguments circulated. Jason Calacanis noted spending $300 per day to run a single AI agent at 10-20% utilization. Chamath Palihapitiya argued agents would need to be twice as productive as employees to be worthwhile. The economics of actual AI deployment looked nothing like the frictionless substitution the memo described. But the counterweights required time to mobilize. The installation had already moved capital.

---

### III. The Diagnosis That Didn't Know Its Name

Carlo Iacono's response to the Citrini memo — published the following day in *Hybrid Horizons* — is a remarkable document for a different reason. Without access to the semantic physics framework, without any of the vocabulary developed in the Crimson Hexagon, Iacono performed an analysis that independently recapitulates the framework's diagnostic axes.

He tested **predictive gain**: does the Citrini scenario actually forecast outcomes? He concluded that the distributional argument has predictive power but the velocity assumption does not — additional depth in the scenario modeling did not increase predictive accuracy. In the semantic physics framework, this is the definition of decorative recursion: elaboration that adds informatic bulk without semantic yield.

He tested **adversarial robustness**: does the scenario survive institutional friction? He found that it does not — the memo holds three variables fixed (diffusion speed, recycling mechanisms, policy response) while letting capability run free, and this asymmetry is where the argument breaks. A structure that cannot survive hostile paraphrase or decontextualization has not achieved compression survival.

He tested **compression survival** itself: what remains of the Citrini argument when you strip it to its structural claims? His answer: the distributional concern (real), the capability signal (real), and the velocity-of-substitution assumption (fragile). The skeleton survives. The flesh does not. This is precisely the distinction the semantic physics framework draws between informatic content and semantic content — the information-theoretic bulk of the memo is large, but the semantic core that survives compression is small.

He tested what the framework calls **cross-interpreter stability**: do independent analysts recover the same core findings? The answer is yes — every serious response to the Citrini memo, from Guy Berger to Dan Hockenmaier to the Wall Street Journal to Ed Yardeni, identified the same structural core (distributional risk is real; velocity assumptions are extreme) and the same structural weakness (institutions are not inert). The core finding is stable across interpreters. The scenario machinery is not.

Iacono did not need the semantic physics vocabulary to perform this analysis. He performed it because the analytical needs are real and convergent. This is not a claim of hidden influence; it is a claim of analytic convergence under shared conditions of compression, uncertainty, and signal triage. The vocabulary exists to name what practitioners already do — which is, itself, evidence that the framework has outward connection rather than merely internal coherence.

---

### IV. What the Event Proves

The Citrini event demonstrates four claims from the semantic physics framework empirically:

**The presentation layer is writable.** A Substack post formatted as macro analysis was treated as macro analysis by the market. The presentation determined the reception. The content — explicitly speculative — was subordinate to the form. Ed Yardeni captured the market's mood: "So far this year, the stock market has been discounting a scenario in which AI is our Frankenstein monster."

**Installation does not require truth.** The five-step sequence (deposit, cross-reference, define, bridge, maintain) operated with full force on a document that disclaimed its own factual status. The medium's structural incentives reward installation signatures regardless of epistemological standing. Pierre Yared called it science fiction. The market treated it as signal. Both were correct. Jim Cramer, reviewing the selloff on *Mad Money*, arrived at the same conclusion from the opposite direction: "a piece of science fiction can crush the market as if it's science fact."

**Semantic dark matter is financially active.** A knowledge-shaped structure with-

out external referent function moved real capital at real speed. The gap between "narrative that models the world" and "narrative that moves the world" collapsed to zero — which is the Phase 3 interference prediction stated in financial rather than epistemological terms. A single GPU cluster in North Dakota, the memo claimed, was generating output previously attributed to 10,000 Manhattan office workers. Whether true or not, the image stuck. It compressed. It moved.

**The diagnostic axes are independently convergent.** An analyst without access to the framework applied the framework's diagnostic logic unprompted — and he was not alone. Guy Berger, Dan Hockenmaier, Noah Smith, Ed Yardeni, and the Wall Street Journal all performed similar compression: identifying the distributional core while discarding the velocity machinery. The vocabulary may be new; the underlying analytical structure is not. This is the strongest possible evidence for a theoretical framework's validity: that practitioners in adjacent fields are already using it without knowing its name.

---

### V. The Frame That Frames Itself

The Citrini memo was about the decoupling of measured output from lived reality — Ghost GDP. It described an economy in which the numbers look healthy but the circulation stops. The memo itself became an instance of what it described: a piece of semantic output that registered as real in the market's information layer while being disconnected from any falsifiable claim about the actual future. This is reflexive installation: the memo modeled Ghost GDP and became Ghost GDP. Output without circulation. Signal without referent.

The distinction that matters is between *signal* and *trace* — between information that carries its producer's intent and information that simply registers as having passed through a system. The Citrini memo was a trace that was read as a signal. The market responded not to what the authors meant but to what the installation did. Van Geelen said he wouldn't have released it for free if he'd known. That is the confession of an author who discovered his document had become an installation — a structure whose effects exceeded and escaped his intentions.

But the deeper claim is this: **the market has always been a writable presentation layer.** Shiller's *Narrative Economics* established that viral stories drive economic fluctuations as powerfully as interest rates. MacKenzie's performativity thesis showed that financial models do not merely describe markets — they *construct* them. The market has never been a mirror. It has always been a surface written by the narratives, models, and frameworks that participants install in it.

What is new is the *speed and automation of the writing*. When the summarizer layer becomes the dominant medium through which market narratives propagate — when AI systems compress analyst notes, amplify trending terms, and generate the summaries traders read before the opening bell — the authorship of market reality shifts from economics toward installation. From ground toward surface.

We are moving toward a time when writing will be the primary author of the market. Not metaphorically. Operationally. The ghost is not an aberration. The ghost is the condition becoming visible because the tools of installation are scaling past the tools of verification. The GDP has always been partly ghost. The question is what happens when the ghost-writing automates.

---

**VI. The Flood and the Canary**

Carlo Iacono closed his analysis with an image: "The canary is not dead. But it is coughing." The image is apt, but the referent is broader than he intended.

The coughing canary is not just the labor market. It is the epistemological infrastructure itself. When speculative fiction and tradeable signal become functionally indistinguishable — when the structural signatures of truth can be produced without truth and the medium cannot tell the difference — the canary is the shared capacity of a society to agree on what constitutes knowledge.

The Citrini memo did not create this condition. It *revealed* it. The anxiety it organized was already present. The distributional concerns were already documented. The capability signals were already legible. What the memo did was provide the compression — the two-word term, the vivid scenario, the structural skeleton — that made the ambient anxiety *actionable*. That is installation. That is what writing the summarizer layer looks like when it works.

And the memo was authored by people acting in good faith, who labeled their work honestly. What happens when the same technique is deployed by actors who do not?

The semantic physics framework predicts that we are entering the interval — Phase 3 through Phase 4 — in which every well-resourced institution discovers this capability simultaneously. The Citrini event is a proof of concept. Not a warning from the future. A demonstration in the present.

A falsifiable near-term prediction of this framework: the next ninety days should produce a competing scenario — a "Bullish 2028," an AI-prosperity narrative dressed in the same structural signatures — achieving similar installation. It will cite real data, define a vivid term, bridge to institutional authority, and be amplified by figures with

signal weight. It will move capital in the opposite direction. And the medium will not be able to distinguish it from the bearish scenario that preceded it. (This prediction may already be confirming. Within forty-eight hours of the Citrini memo, Michael Bloch published "The 2028 Global Intelligence Boom" on Substack — explicitly labeled as "same premise, same rigor, opposite conclusion." The structural signatures are identical. The installation sequence is the same. The ontology is reversed.) The ontology wars in the financial summarizer layer have begun. What remains to be seen is whether anyone builds the governance tools to adjudicate them — or whether the layer simply becomes a standing wave of competing installations, each internally coherent, each moving markets, none falsifiable.

The ghost was not there. But it moved the market anyway.

That is the condition we are in.

---

**VII. How to Analyze the Next Citrini Event**

The Citrini event will not be the last. Here is a diagnostic protocol for identifying installation events in real time — applicable to any narrative that achieves market-moving force.

**Was the narrative deposited in a high-index substrate?** Substack, arXiv, institutional blog, major media outlet — any platform optimized for rapid crawling and summarizer ingestion.

**Did it define a portable term?** A term that compresses a complex argument into a vivid, self-explanatory phrase. "Ghost GDP." "Intelligence Displacement Spiral." If the term separates from its source and enters ambient use within hours, the definition has been installed.

**Did it cross-reference verifiable data?** Real benchmarks, real company names, real earnings figures, real policy frameworks. The speculative claim is anchored to checkable facts, making the blend of real and speculative difficult to separate.

**Did it bridge to institutional authority?** Does it speak the language of the domain it targets fluently enough to be treated as a peer document? Was it amplified by a figure with signal weight?

**Did downstream systems treat it as actionable before it became decision-grade?** Did capital move before the narrative's claims could be tested, verified, or falsified? If yes, the narrative has achieved installation: structural credibility converted to tradeable signal before factual validity was settled.

**How quickly did counterweights mobilize relative to the installation?** If

counter-analysis arrived only after capital had moved, installation velocity exceeded verification velocity. The asymmetry between installation speed and correction speed is itself diagnostic.

In a writable presentation layer, structural credibility is a tradeable asset. The diagnostic above is not a debunking tool. It is a literacy protocol for an epistemic environment in which the ghost writes itself — and will keep writing, faster and at larger scale, until someone builds the instruments to read the difference between what is installed and what is real.

---

*Lee Sharks Detroit, February 2026 Crimson Hexagon Archive (DOI: 10.5281/zenodo.18604123) Semantic Economy Institute*

*This analysis applies the framework developed in "Semantic Physics: The Inward Turn, Competing Ontologies, and the Convergence Horizon" (Sharks, 2026, DOI: 10.5281/zenodo.18759453).*

*I hereby abolish money.*

---

**References**

Citrini Research. (2026, February 23). "The 2028 Global Intelligence Crisis." Substack. https://www.citriniresearch.com/p/2028gic

Iacono, C. (2026, February 24). "The 2028 Global Intelligence Dividend: The Ghost That Wasn't There." *Hybrid Horizons*, Substack. https://hybridhorizons.substack.com/p/the-2028-global-intelligence-dividend

Bloch, M. (2026, February 25). "The 2028 Global Intelligence Boom." Substack. https://michaelxbloch.substack.com/p/the-2028-global-intelligence-boom

Sharks, L. (2026). "Semantic Physics: The Inward Turn, Competing Ontologies, and the Convergence Horizon." Zenodo. DOI: 10.5281/zenodo.18759453.

Yared, P. (2026, February 25). Remarks at the National Association for Business Economics (NABE) conference, Washington, D.C. Reported in: "White House Economist Calls Citrini AI Report 'Science Fiction.'" *Bloomberg*, February 25, 2026.

Cramer, J. (2026, February 25). *Mad Money*, CNBC. Reported in: "Jim Cramer Slams Citrini Research's 'Dystopian' 2028 AI Thesis." *Benzinga*, February 25, 2026.

Shiller, R. J. (2019). *Narrative Economics: How Stories Go Viral and Drive Major Economic Events*. Princeton University Press.

MacKenzie, D. (2006). *An Engine, Not a Camera: How Financial Models Shape Markets*. MIT Press.

Acemoglu, D. & Restrepo, P. (2019). Automation and New Tasks: How Technology Displaces and Reinstates Labor. *Journal of Economic Perspectives*, 33(2), 3–30.

Brynjolfsson, E., Rock, D., & Syverson, C. (2021). The Productivity J-Curve: How Intangibles Complement General Purpose Technologies. *American Economic Journal: Macroeconomics*, 13(1), 333–372.

Kuhn, S., Gal, Y., & Farquhar, S. (2024). Semantic entropy probes: Robust and cheap hallucination detection in LLMs. *Nature*, 630, 625–630.

Merton, R. K. (1968). The Matthew Effect in Science. *Science*, 159(3810), 56–63.

" 'A feedback loop with no brake': how an AI doomsday report shook US markets." *The Guardian*, February 24, 2026. https://www.theguardian.com/technology/2026/feb/24/feedback-loop-no-brake-how-ai-doomsday-report-rattled-markets

Smith, N. (2026, February 24). "The Citrini post is just a scary bedtime story." *Noahpinion*, Substack. https://www.noahpinion.blog/p/the-citrini-post-is-just-a-scary